

Decomposing the Annuity Puzzle: Rational, Preference, and Behavioral Channels under a UK-Anchored Wedge Transport*

Derek Tharp[†]

May 11, 2026

Abstract

Yaari (1965) showed that a consumer with uncertain lifetime and no bequest motive should fully annuitize under actuarially fair pricing. Observed voluntary annuity ownership among US retirees is only a few percent. I build a lifecycle model that begins from a standard rational and preference-based account of annuity demand and then adds a *bundled behavioral wedge* representing three conceptually distinct phenomena that activate when compulsion lifts: source-dependent utility (SDU; Blanchett and Finke, 2024, 2025); narrow-framing purchase-event disutility (PED; Barberis and Huang, 2009; Tversky and Kahneman, 1992); and choice-architecture salience (provider-communication, adviser-conversation, status-quo, and default-vs-opt-in framing effects; Madrian and Shea, 2001; Brown et al., 2008). The bundle is identified externally from a single market-design moment—the UK 2015 pension-freedoms reform—and applied to the model’s no-behavioral baseline as a deterministic

*Generative AI tools were used to assist with code development, manuscript preparation, and editing. All analysis, interpretation, and conclusions are the sole responsibility of the author. Replication code and data are available at <https://github.com/DerekTharp/annuity-puzzle-model>.

[†]Department of Accounting & Finance, University of Southern Maine. Email: derek.tharp@maine.edu.

multiplicative transformation: predicted US ownership equals the no-behavioral baseline multiplied by the UK proportional retention factor (post/pre $\approx 17/95 = 0.18$ at the production midpoint). The three named components are not separately parameterized in the model; they are observed jointly in UK 2015 and transported jointly. Without the wedge the model substantially overshoots observed US ownership; applying the wedge to the no-behavioral baseline yields the production prediction at the production midpoint, with a sensitivity range across UK retention anchors that contains both HRS empirical measures.

The behavioral bundle is calibrated under a *proportional behavioral wedge* transport: the ratio of post-reform UK voluntary retention to pre-reform UK compulsory retention (approximately $17/95 \approx 0.18$ at the production midpoint) is treated as country-invariant and applied to the model’s pre-behavioral US baseline (33.6%) to yield the production target of approximately 6%. The 2015 reform shifted the regulated regime for DC pension pots from de facto compulsory annuitization—enforced pre-reform by a 55% unauthorised-payment charge on alternative withdrawals—to voluntary access with marginal-rate tax neutrality, producing post-reform voluntary retention of 13–25% across ABI/FCA estimates. The empirical fact that annuity share remained at approximately 9% in 2024/25 despite gilt yields at 14-year highs ([Financial Conduct Authority, 2025](#)) indicates that the post-2015 retention regime has been broadly stable across a full rate cycle, ruling out a substantial rate-environment confound. The UK retention range maps to a US ownership bracket of approximately [4.6%, 8.8%], with a wider sensitivity range across alternative anchors of approximately [2.5%, 8.8%]. No US moment is used to discipline the behavioral bundle; consistency with US ownership is an out-of-sample comparison.

The bundled wedge is implemented as multiplicative post-processing on ownership rates rather than as a structural parameter inside the Bellman equation: the model produces the no-behavioral baseline (rational + preferences + the structural public-care aversion channel χ_{LTC}), and the wedge is applied externally as a deterministic scalar reduction. The Blanchett–Finke spending differential is descriptive evidence that

source-dependent utility exists in retirement spending data, but no separate identifying moment is used; the bundle is identified jointly from UK 2015 alone. Counterfactuals that “remove the bundled wedge” are interpreted as removing the multiplicative reduction, not as removing a single named force. Channel attribution among the model-internal channels uses exact Shapley values computed over all $2^9 = 512$ subsets of the nine model-internal channels; the bundled wedge appears separately as a multiplicative reduction, not as a Shapley channel.

JEL Codes: D14, D15, D91, G22, H55, J14

Keywords: annuity puzzle, lifecycle model, Shapley decomposition, stochastic mortality, age-varying consumption needs

1 Introduction

Yaari (1965) established one of the sharpest predictions in consumer theory: an individual facing uncertain lifetime, with no bequest motive and access to actuarially fair annuities, should convert all wealth into a life annuity. The mortality credit—the return premium financed by redistributing wealth from decedents to survivors—dominates any asset with the same underlying return. Davidoff et al. (2005) extended this result to incomplete markets, showing that partial annuitization remains optimal under a wide range of conditions. Yet observed voluntary annuity ownership among single US retirees aged 65–69 is only a few percent: 2.02% on the cleaner HRS lifetime annuity contract indicator (q286 series) and 3.34% on the conventional any-annuity income proxy used in prior literature (Lockwood, 2012; Dushi and Webb, 2004; Poterba et al., 2011).¹ Six decades of research have proposed partial explanations for this gap, but no single paper has nested all the major channels in a unified framework and measured their joint quantitative effect.

This paper provides a calibrated structural accounting exercise for that gap. I begin from a standard rational and preference-based account of annuity demand and then add two behavioral channels grounded in independent empirical evidence. The standard account comprises six *rational* channels (pre-existing Social Security annuitization, bequest motives, combined medical expenditure risk and the health-mortality correlation identified by Reichling and Smetters, 2015, subjective survival pessimism (O’Dea and Sturrock, 2023), pricing loads, and inflation erosion) and two *preference* channels (age-varying consumption needs (Aguiar and Hurst, 2013) and state-dependent utility (Finkelstein et al., 2013)). On its own, this eight-channel account substantially overshoots observed US ownership.

The paper then adds a *bundled behavioral wedge* representing three named phenomena that activate when compulsion lifts. Source-dependent utility (SDU) captures the post-

¹The two HRS measures differ because the income proxy (*rwinn*) classifies any positive annuity income as ownership, including one-time DC pension withdrawals and short-period payouts; the lifetime contract indicator (the question stem “Will this continue for the rest of your life?”) isolates true life-contingent annuity contracts. Section 3 discusses both measures in detail.

purchase “license to spend” that retirees apply to guaranteed income flows but not portfolio assets (Shefrin and Thaler, 1988; Thaler, 1999; Blanchett and Finke, 2024, 2025). Narrow-framing purchase-event disutility (PED) captures the at-purchase loss aversion over the unrecouped premium until breakeven (Tversky and Kahneman, 1992; Barberis and Huang, 2009; Brown et al., 2008; Chalmers and Reuter, 2012). Choice-architecture salience captures the sub-bundle of provider-communication framing, adviser-conversation effects, status-quo bias, and Madrian–Beshears default-vs-opt-in inertia (Madrian and Shea, 2001; Goda et al., 2014). The bundle is identified from a single external moment—the UK 2015 pension-freedoms reform—under a *proportional behavioral wedge* transport that treats the post/pre retention ratio as country-invariant. The wedge is applied to the model’s no-behavioral baseline as a deterministic multiplicative reduction; no behavioral parameter is calibrated inside the Bellman equation. The 2015 reform bundled tax, advice, default, and product-market changes; the empirical evidence is best read as a market-design anchor rather than a clean causal estimate of any single behavioral primitive in isolation. The English Longitudinal Study of Ageing (ELSA) pension grid provides supporting individual-level evidence: 90.2% of wave 6 (2012–13) DC pension recipients reported annuity-style income under the pre-freedoms mandate, and pooling waves 8–11 (2016–24), 1.27% of observed lump-sum disposition records used the lump sum to buy annuity income (n=869). No US moment is used to discipline the bundle; consistency with US ownership is an out-of-sample comparison.

The paper’s contribution is fourfold. First, the unified framework—the first to nest age-varying consumption needs alongside the established rational and preference channels of the annuity-puzzle literature. Second, the formal taxonomy of three named behavioral phenomena (SDU, PED, choice-architecture salience) treated as a single bundled wedge identified from a single external moment. Third, the proportional-wedge transport from the UK 2015 pension-freedoms reform as a disciplined identifying assumption, with explicit cross-country preference-homogeneity assumption. Fourth, an exact Shapley decomposition over all $2^9 = 512$ subsets of the nine model-internal channels that attributes the demand

reduction without order dependence.

The results have a three-layer structure. First, under modern annuity pricing ($MWR = 0.87$, [Wettstein et al., 2021](#)), the six rational channels alone predict ownership of 8.5% relative to a frictionless population benchmark of 41.8%; adding the two preference channels brings the prediction to 1.8%. The standard rational and preference account therefore substantially overshoots the empirical targets—echoing the residual gap in [Pashchenko \(2013\)](#) and [Peijnenburg et al. \(2016\)](#). Second, adding Force A alone (source-dependent utility) raises predicted ownership to 27.8% by unlocking the post-purchase “license to spend” on annuity income—confirming that one component of the behavioral bundle in isolation moves predictions in the wrong direction. Third, adding Force B with the production normalization yields the bundled equilibrium prediction of approximately 6% predicted US voluntary ownership, with a sensitivity range across UK retention anchors of approximately [4.6%, 8.8%] and a wider range across all anchors of approximately [2.5%, 8.8%]. Both HRS measures of US lifetime annuity ownership—the conventional any-annuity income proxy (3.34%, 95% CI [2.89%, 3.85%]) and the cleaner lifetime-contract indicator (2.02%, 95% CI [1.68%, 2.43%])—fall within the wider sensitivity range. The bundle is identified entirely from non-US data; consistency with US ownership is an out-of-sample comparison.

Relative to prior multi-channel papers, the present model broadens the channel set and changes the attribution. [Pashchenko \(2013\)](#) jointly considered Social Security, bequests, minimum purchase requirements, and pricing frictions, but not the Reichling–Smetters health-survival mechanism, subjective survival pessimism, or age-varying needs. [Peijnenburg et al. \(2016\)](#) emphasized incomplete markets and background risk, concluding that the annuity puzzle remains unresolved in their framework. This paper shows that pricing, correlated health and mortality risk, and late-life expenditure decline are quantitatively central once evaluated in a common model.

The exact Shapley decomposition over the nine model-internal channels is the paper’s preferred attribution result for the no-behavioral baseline. Among demand-suppressing model

channels, pricing loads (14.0 pp, 33%) and the combined medical risk and Reichling–Smetters health-mortality correlation channel (8.6 pp, 21%) are the two largest contributors, followed by age-varying consumption needs (2.9 pp, 7%), survival pessimism (5.9 pp, 14%), and inflation erosion (1.4 pp, 3%). Social Security enters with a negative Shapley value (−17.6 pp, −42%), reflecting its role as an ownership booster. The bundled behavioral wedge does not appear in the Shapley enumeration: it is applied as multiplicative post-processing on the ownership rate produced by the no-behavioral baseline. State-dependent utility and bequest motives both contribute negligibly (0.3 pp and 0.9 pp respectively), entering the paper as completeness checks rather than main drivers.

The policy and welfare results extend beyond the ownership match. Group annuity pricing ($MWR = 0.90$) raises predicted ownership relative to the production calibration. Removing the bundled behavioral wedge (an illustrative default-architecture lever) recovers the no-behavioral baseline of 1.7%—a counterfactual roughly consistent with the UK 2015 reform and the [Goda et al. \(2014\)](#) federal Thrift Savings Plan natural experiment, but reported with the explicit caveat that the wedge is identified only as a bundle and cannot be decomposed into the contributions of its named components. Social Security cut counterfactuals reveal that a 23% benefit cut—the projected consequence of trust fund exhaustion around 2033—raises private annuity demand substantially, but the largest cuts eventually reduce demand by removing the income floor that makes private annuitization feasible at all. Under current pricing the welfare gain from annuity market access is small across the wealth distribution because most calibrated households optimally do not annuitize, so access is approximately neutral; policy interventions that change the optimal choice (group pricing, demand-side architecture) generate the larger welfare gains.

These results shift the question away from whether retirees irrationally fail to annuitize and toward a more policy-relevant one: which channels suppress annuity demand most, and which market reforms would make annuitization welfare-improving for particular households?

Section 2 specifies the model. Section 3 describes the calibration. Section 4 presents the

decomposition, counterfactuals, and welfare analysis. Section 5 reports sensitivity checks. Section 6 discusses limitations. Section 7 concludes.

2 Model

2.1 Environment

Time is discrete. An individual enters at age 65 (period $t = 1$) and may survive to a maximum age of 110 (period $T = 46$). At $t = 1$, the individual makes a single irreversible decision: what fraction $\alpha \in [0, 1]$ of initial liquid wealth W_0 to convert into a single premium immediate annuity (SPIA). The remaining wealth $(1 - \alpha)W_0$ stays liquid. In each subsequent period, the individual chooses consumption c_t from available resources.

The state at each period t is characterized by four variables: liquid wealth $W_t \geq 0$, annuity income A (fixed after the age-65 purchase), health status $H_t \in \{1, 2, 3\}$ (Good, Fair, Poor), and age t . Social Security income $SS(t)$ is received exogenously in each period and is not a decision variable.

2.2 Preferences

The individual maximizes expected discounted lifetime utility. Flow utility over consumption takes the constant relative risk aversion (CRRA) form:

$$u(c) = \frac{c^{1-\gamma}}{1-\gamma}, \quad \gamma > 0, \quad \gamma \neq 1. \quad (1)$$

The coefficient of relative risk aversion γ governs both risk aversion and the elasticity of intertemporal substitution.

Bequest utility follows the De Nardi–French–Jones specification, which nests bequests as a luxury good:

$$v(b) = \theta \frac{(b + \kappa)^{1-\gamma}}{1-\gamma}, \quad (2)$$

where b is bequeathable wealth (liquid wealth at death; annuity income ceases), $\theta \geq 0$ controls bequest intensity, and $\kappa > 0$ is a shifter that determines the degree to which bequests are a luxury good (De Nardi, 2004). When κ is large relative to typical wealth levels, the marginal utility of bequests is nearly flat for low-wealth individuals and declines steeply for the wealthy. This specification, estimated by Lockwood (2012) using Health and Retirement Study (HRS) data, captures the empirical pattern that bequest motives are concentrated among high-wealth households.

The parameters θ and κ must be interpreted jointly. With the DFJ luxury-good calibration ($\theta = 56.96$, $\kappa = \$272,628$), an individual with \$50,000 in wealth faces negligible marginal bequest utility, while an individual with \$1,000,000 faces strong resistance to annuitizing. Combining this θ with a small κ (e.g., \$10, the homothetic specification) produces pathological results: extreme marginal bequest utility near $b = 0$ acts as a precautionary buffer rather than a genuine bequest motive, anomalously *increasing* predicted ownership above the no-bequest case.² The DFJ estimates from Lockwood (2012) are jointly calibrated to HRS bequest data; using one parameter with the other replaced is not a valid counterfactual.

The bequest channel’s isolated effect on ownership is modest under the DFJ specification (100% retention rate in the sequential decomposition) because the luxury-good κ makes the bequest motive nearly irrelevant for the median retiree in the HRS sample. The 100% retention rate in the sequential decomposition reflects the ordering: bequests are added when ownership is already 100% (steps 0–1 include only SS). In the full model, removing bequests raises ownership from 1.7% to 21.0% (Table 8), because bequests interact with pricing loads and inflation to suppress demand for agents near the annuitization margin. The Shapley decomposition, which averages over all orderings, assigns bequests a value of 0.9 pp (2% share).

²At $\kappa = \$10$, marginal bequest utility near $b = 0$ is $\theta \cdot 10^{-\gamma}$, which is extreme relative to marginal consumption utility. Individuals avoid annuitizing not to protect bequests, but to avoid near-zero-wealth states where $v'(b)$ becomes very large. The implementation applies a small numerical floor at $\max(b + \kappa, 1)$ to prevent overflow at $b = 0$ when $\kappa = 0$ is used as a robustness check; with the production DFJ calibration ($\kappa = \$272,628$) the floor never binds. See Appendix C for details.

The individual discounts future utility at rate β per period.

2.3 Health dynamics and mortality

Health follows a first-order Markov process with three states: Good ($H = 1$), Fair ($H = 2$), and Poor ($H = 3$). Transition probabilities are age-dependent and calibrated to HRS panel data following [De Nardi et al. \(2010\)](#):

$$\Pr(H_{t+1} = h' \mid H_t = h, \text{age} = t) = \pi(h, h', t). \quad (3)$$

Mortality depends on health through a proportional hazard specification. Baseline survival probabilities $s_{\text{base}}(t)$ come from the SSA period life table used by [Lockwood \(2012\)](#). Health adjusts survival through:

$$s(t, H) = s_{\text{base}}(t)^{\mu(H, t)}, \quad (4)$$

where $\mu(H, t)$ is a health- and age-specific hazard multiplier with $\mu(\text{Good}, t) < 1 < \mu(\text{Poor}, t)$ and $\mu(\text{Fair}, t) = 1$. This specification maps directly to proportional hazards: if the baseline hazard is $\lambda_{\text{base}}(t) = -\ln s_{\text{base}}(t)$, then the adjusted hazard is $\mu(H, t) \cdot \lambda_{\text{base}}(t)$.

The baseline specification uses constant hazard multipliers $\mu = [0.50, 1.0, 3.8]$ for Good, Fair, and Poor health respectively, chosen as a central tendency across age bands from the HRS data. The health-mortality gradient compresses with age: the Poor-to-Fair hazard ratio falls from 3.29 at ages 65–74 to 2.77 at 75–84 to 1.82 at 85+. [Section 5](#) reports results across the empirical range of multiplier specifications.

The age-varying multipliers are the mechanism underlying the [Reichling and Smetters \(2015\)](#) result. A negative health shock simultaneously increases mortality ($\mu(\text{Poor}, t) > 1$ reduces expected remaining annuity payments) and, through the medical expenditure process described below, increases the need for liquid wealth. The annuity becomes “valuation-risky”: it loses value in precisely the states where liquidity is most valuable. The effect is strongest at

younger ages where the gradient is widest, which matters because the annuitization decision is made at age 65.

2.4 Subjective survival beliefs

O’Dea and Sturrock (2023) document that individuals aged 65–69 underestimate their survival probability at age 75 by approximately 15 percentage points relative to actuarial tables (subjective $\Pr(\text{survive to } 75) \approx 0.71$ versus actuarial 0.86). This pessimism reduces the perceived value of annuities by overweighting the probability of dying before recovering the premium.

The model allows the agent’s subjective survival probabilities to differ from the objective rates used by the insurer to price annuities:

$$s^{\text{subj}}(t, H) = \psi \cdot s(t, H), \quad \psi \in (0, 1]. \quad (5)$$

The parameter ψ scales each one-year survival probability downward. The annuity is priced using the objective survival rates $s(t, H)$; only the agent’s Bellman equation uses s^{subj} . A per-year factor of $\psi = 0.960$ matches the O’Dea–Sturrock finding: over a 10-year horizon, $(0.71/0.86)^{1/10} \approx 0.960$. When $\psi = 1$, beliefs are objective and this channel is inactive.

2.5 Medical expenditures

Out-of-pocket medical expenditures depend on age and health. Log medical spending follows:

$$\ln m_t = \mu_m(t, H_t) + \sigma_m(t, H_t) \cdot \varepsilon_t, \quad \varepsilon_t \sim N(0, 1), \quad (6)$$

where μ_m and σ_m are the mean and standard deviation of log expenditures, both age- and health-dependent. The baseline calibration matches the age profiles reported by Jones et al. (2018): mean out-of-pocket spending of \$4,200 at age 70, rising to \$29,700 at age 100, with

a 95th percentile of approximately \$111,200 at age 100.

A Medicaid safety net provides a consumption floor $\underline{c}$. If total resources (wealth plus income minus medical costs) fall below $\underline{c}$, the government covers the shortfall. The individual consumes $\underline{c}$ and enters the next period with zero liquid wealth. This floor is calibrated to the SSI federal benefit rate plus average state supplement for an elderly individual (\$6,180, following the value used in [Lockwood, 2012](#)).

Expectations over medical expenditure shocks are computed using nine-node Gauss–Hermite quadrature. At the calibrated variance ($\sigma \approx 1.4$), lower-order rules place insufficient weight in the tails of the lognormal distribution; Appendix [G](#) documents the convergence diagnostics.

2.6 Annuity pricing

The annuity is nominal: the insurer prices the stream of nominal payments using the nominal discount rate $r_{\text{nom}} = (1+r)(1+\pi) - 1$, where r is the real interest rate and π is the inflation rate. The actuarially fair annual payout per dollar of premium is:

$$\text{payout}_{\text{fair}} = \frac{1}{\sum_{k=0}^{T-1} \frac{\prod_{j=0}^{k-1} s_{\text{base}}(j)}{(1+r_{\text{nom}})^k}}. \quad (7)$$

The nominal discount rate reflects the insurer’s investment in nominal bonds. The resulting initial payout is higher than under real discounting, but the consumer’s real income from the annuity declines each period:

$$A_t^{\text{real}} = \frac{A_1}{(1+\pi)^{t-1}}, \quad (8)$$

where π is the annual inflation rate. At $\pi = 2\%$, purchasing power falls by 39% over 25 years.

When $\pi = 0$ (the Yaari benchmark and intermediate decomposition steps where inflation

is inactive), $r_{\text{nom}} = r$ and the pricing reduces to the standard real-annuity formula.

The annuity is priced with a money’s worth ratio (MWR), defined as the expected present value of payouts per dollar of premium. The loaded payout rate is $\text{MWR} \times \text{payout}_{\text{fair}}$. I set $\text{MWR} = 0.87$ at baseline, consistent with the population-table estimates of [Mitchell et al. \(1999\)](#) for the US individual annuity market. This choice merits discussion. [Mitchell et al. \(1999\)](#) estimated MWRs of 0.80–0.85 using population mortality tables, which is the relevant benchmark for an individual considering annuitization (the “money’s worth” of the purchase against population life expectancy). Using annuitant mortality tables—which adjust for selection—yields higher MWRs of 0.90–0.94, but this overstates the value proposition for the marginal non-annuitant. More recent estimates from [Wettstein et al. \(2021\)](#) suggest some improvement, with population-table MWRs around 0.85. Section 5 reports full sensitivity: at $\text{MWR} = 0.85$, predicted ownership rises to 0.4%; at $\text{MWR} = 0.90$, to 5.8%. The steep sensitivity to MWR is itself a finding—it identifies pricing as the most actionable policy lever.

A fixed cost of \$2,500 is incurred upon any purchase.³

2.7 Budget constraint

Within each period, the timing is:

1. The individual enters with liquid wealth W_t , annuity income A_t , Social Security $\text{SS}(t)$, and health H_t .
2. Medical expenditure m_t is realized.
3. The individual chooses consumption c_t .
4. Remaining wealth earns the risk-free return r .

³The MWR is computed using beginning-of-period payment timing: the first payment occurs at purchase. This convention is standard in the annuity pricing literature ([Mitchell et al., 1999](#)) and matches the implementation in the replication code.

5. Health transitions to H_{t+1} .
6. The individual survives to $t + 1$ with probability $s(t, H_t)$.

The intertemporal budget constraint is:

$$W_{t+1} = (1 + r)(W_t + A_t + SS(t) - m_t - c_t), \quad (9)$$

subject to $c_t \geq \underline{c}$ and $W_{t+1} \geq 0$ (no borrowing).

2.8 Age-varying consumption needs

[Aguiar and Hurst \(2013\)](#) document that consumption expenditure declines at roughly 2% per year in retirement, even among wealthy households, once time costs and home production are accounted for. This decline reflects changing consumption needs—reduced transportation, clothing, and food-away-from-home expenditures—rather than binding liquidity constraints. The declining profile reduces the value of late-life income streams, including annuities.

The model captures this through an age-dependent weight on flow utility:

$$w(t) = (1 - \delta_c)^{t-1}, \quad \delta_c = 0.02, \quad (10)$$

where $t = 1$ at age 65. Flow utility becomes $w(t) \cdot u(c)$. This specification is equivalent to an age-dependent felicity shifter: as consumption needs decline, the marginal value of additional consumption in late life falls, reducing the insurance value of an annuity that provides a level real income stream. When $\delta_c = 0$, this channel is inactive and the model reduces to the standard time-separable case.

Age-varying consumption needs should be distinguished from time preference (β). The discount factor β governs the rate at which future utility is discounted; $w(t)$ governs how much utility a given level of consumption generates at each age. A constant β with declining $w(t)$ implies that the individual values future consumption opportunities but needs less

consumption to achieve the same felicity as she ages.

2.9 State-dependent utility

[Finkelstein et al. \(2013\)](#) estimate that the marginal utility of consumption declines with poor health, with a central estimate implying that individuals in poor health value a dollar of consumption roughly 10–25% less than those in good health. [Reichling and Smetters \(2015\)](#) incorporated this channel in their annuity model. State-dependent utility reduces annuity demand because the annuity delivers income in poor-health states where its marginal utility is lower.

The model allows the marginal utility of consumption to depend on health status through a multiplicative weight:

$$u(c, H) = \varphi(H) \cdot \frac{c^{1-\gamma}}{1-\gamma}, \quad (11)$$

where $\varphi(\text{Good}) = 1.00$, $\varphi(\text{Fair}) = 0.92$, and $\varphi(\text{Poor}) = 0.82$, following the mapping used by [Reichling and Smetters \(2015\)](#) based on the estimates of [Finkelstein et al. \(2013\)](#). When $\varphi(H) = 1$ for all H , this channel is inactive.

In practice, state-dependent utility has a negligible effect on predicted ownership (Shapley value of 0.3 pp). The channel is included for completeness and to confirm that its quantitative irrelevance, suggested by the prior literature, holds in the unified framework.

2.10 The bundled behavioral wedge

The nine rational and preference-based channels omit a set of behavioral phenomena that activate when retirees face a free annuitization choice. The empirical literature documents at least three such phenomena operating at distinct decision moments. **Force A: source-dependent utility** ([Shefrin and Thaler, 1988](#); [Thaler, 1999](#); [Blanchett and Finke, 2024, 2025](#)) operates *post*-purchase and raises annuitization by elevating the consumption rate retirees apply to guaranteed income flows relative to portfolio assets. **Force B: narrow-**

framing purchase-event disutility (Tversky and Kahneman, 1992; Barberis and Huang, 2009) operates *at*-purchase and suppresses annuitization through loss aversion over the unrecouped premium until breakeven. **Force C: choice-architecture salience** (Madrian and Shea, 2001; Brown et al., 2008; Chalmers and Reuter, 2012) operates *around*-purchase and is itself a sub-bundle of provider-communication framing, adviser-conversation effects, status-quo bias, and the Madrian–Beshears family of default-vs-opt-in inertia channels; the empirical literature documents these effects at the 5–10 percentage-point order of magnitude in retirement decision settings, smaller than Force A or Force B but not negligible, and with mixed sign across sub-components.

I treat all three forces as components of a single *bundled behavioral wedge* that is jointly identified from one external moment—the UK 2015 pension-freedoms reform. The reform turned Force A, Force B, and Force C on simultaneously when compulsion lifted: source-dependence, narrow framing, and choice-architecture salience all became binding at the same moment when retirees were given a free choice. The resulting attrition is the joint footprint of the three forces, which is what the data permit transporting.

Mechanism: multiplicative transport. The bundled wedge is applied as a deterministic multiplicative transformation to the model’s no-behavioral baseline. Under proportional-wedge identification, the ratio of post-reform UK voluntary retention to pre-reform UK compulsory retention is treated as country-invariant:

$$\text{predicted US ownership} = \text{no-behavioral US baseline} \times \frac{\text{UK retention}_{\text{post}}}{\text{UK retention}_{\text{pre}}}. \quad (12)$$

Forces A, B, and C are NOT parameterized in the model itself. They are named in the manuscript as conceptually distinct components of the bundle (Section 2.10), but no single force is separately mechanized. The model’s role is to produce the no-behavioral baseline (rational + preferences + structural channels including χ_{LTC}); the wedge is applied as post-processing in scripts/export_manuscript_numbers.jl.

The proportional retention factor at the production midpoint is $17/95 \approx 0.18$. The UK retention sensitivity range of [13%, 25%] maps to factors of [0.137, 0.263]. Applied to the no-behavioral baseline of 1.7%, the wedge produces a predicted US ownership bracket of [5.7%, 11.0%] with production midpoint 7.5%.

Why mechanism-free. The model could in principle parameterize each behavioral channel structurally—an SDU consumption transformation $c_{\text{eff}} = c_{\text{income}} + \lambda_W c_{\text{portfolio}}$ for Force A, an at-purchase loss-aversion penalty ψ_{purchase} for Force B—and calibrate them via SMM bisection. We considered and rejected this approach for two reasons. First, two parameters cannot be separately identified from a single bundling moment without auxiliary anchors that re-introduce the multi-moment structure that bundled identification rejects. Second, the at-purchase penalty mechanism produces non-smooth participation thresholds at low pre-behavioral baselines (where χ_{LTC} leaves little room for the wedge to operate); the threshold behavior does not faithfully implement the multiplicative transport interpretation. Multiplicative transport as post-processing is the cleanest implementation of proportional-wedge identification.

Cross-country transport: maintained assumption. Proportional-wedge identification is a transport assumption, not a property of the data. The maintained assumption is that the equilibrium retention ratio (post-reform / pre-reform) is country-invariant: UK and US retirees have sufficiently similar preference distributions over the SDU, PED, and choice-architecture-salience phenomena that the same proportional reduction applies in both populations. This assumption is load-bearing. Specifically: (i) the cross-country preference homogeneity of the underlying behavioral phenomena, (ii) sample exchangeability between UK 2015 DC pension-pot decumulators and US 65–69 retirees in the HRS sample, and (iii) institutional comparability sufficient that the same retention ratio applies. UK and US institutional environments differ along several dimensions not absorbed by the wedge: DC pension-pot scale and concentration, the FCA/RDR adviser regime versus US insurance-

distribution structures, Medicaid/SS interaction patterns, and minimum-purchase requirements. The wedge bracket [5.7%, 11.0%] is therefore an out-of-sample extrapolation under a transport assumption, not a clean identification result.

Tax-penalty and rate-environment confounds. Two confounds within the UK identifying moment itself are worth flagging. First, the UK 55% unauthorised-payment charge that enforced pre-reform compulsion is a budget-constraint mechanism not a preference parameter; the proportional-wedge identification absorbs this into the bundle but the implied “behavioral” attrition therefore includes a UK-specific tax-removal response (estimated at 5–15 pp of the raw 75–87 pp swing; see [Cribb and Emmerson, 2018](#); [Cannon et al., 2016](#) for analyses of the UK reform’s tax and behavioral components). Second, the rate-environment confound is empirically small: FCA Retirement Income Market Data ([Financial Conduct Authority, 2025](#)) report annuity share of DC pots accessed at approximately 9% in 2024/25 despite gilt yields at 14-year highs, indicating that the post-2015 retention regime has been broadly stable across a full rate cycle and that gilt-yield variation explains 0–5 pp at most of the post-reform retention level. The wider sensitivity range covers both confounds.

Out-of-sample comparison. The calibration is set entirely from UK data, with no reference to observed US ownership. The predicted US voluntary annuitization bracket of [5.7%, 11.0%] is therefore an out-of-sample comparison. Both HRS measures of US lifetime annuity ownership (2.02% and 3.34%) are reported alongside the prediction in [Section 4](#).

Conceptually distinct margins, jointly identified. Within the bundle, the three named forces target empirically distinct margins. Source-dependent utility (Force A) operates on the *ongoing consumption* from existing wealth (post-purchase, every period); it shows up in passive spending patterns documented by Blanchett–Finke and in the fungibility violations documented by [Hastings and Shapiro \(2013\)](#). The narrow-framing penalty (Force B) operates on the *at-purchase decision* (one-time, with stream-form ongoing pain until breakeven); it shows

up at transaction boundaries documented by Brown et al. and in the cumulative-prospect-theory mechanism formalized by [Hu and Scott \(2007\)](#). Choice-architecture salience (Force C) operates on the *decision environment surrounding the purchase* (provider communications, adviser scripts, default-vs-opt-in framing); it shows up in the UK 2015 quasi-experiment, the TSP defaults ([Goda et al., 2014](#)), and the broader Madrian–Beshears defaults literature ([Madrian and Shea, 2001](#)). The literatures are distinct, but the UK identifying moment activated all three at once and cannot, on its own, attribute the swing to any one of them. Sub-decomposing the wedge into Force A / Force B / Force C contributions is therefore not pursued.

2.11 Bellman equation

For $t < T$, the value function satisfies:

$$\begin{aligned} V(W, A, H, t) = \max_c \bigg\{ & w(t) \cdot u(c, H) \\ & + \beta s(t, H) \mathbb{E}[V(W', A', H', t+1) \mid H] \\ & + \beta (1 - s(t, H)) v(W') \bigg\}, \end{aligned} \quad (13)$$

where $W' = (1+r)(W + A_t + \text{SS}(t) - m - c)$, $A' = A/(1+\pi)$ is next period's real annuity income, and the expectation is taken over the joint distribution of H' (health transition) and m (medical shock). The term $w(t)$ is the age-varying needs weight from equation (10) and $u(c, H)$ is the state-dependent utility from equation (11). When both channels are inactive ($\delta_c = 0$, $\varphi(H) = 1$), this reduces to the standard Bellman equation. At the terminal period T , the individual consumes all remaining resources and leaves any residual as a bequest.

At $t = 1$ (age 65), the annuitization decision is:

$$\alpha^* = \arg \max_{\alpha \in [0,1]} V\left((1-\alpha)W_0 - F \cdot \mathbf{1}(\alpha > 0), \alpha W_0 \cdot \text{payout}, H_0, 1\right), \quad (14)$$

where F is the fixed purchase cost.

2.12 Analytical characterization

The numerical decomposition in Section 4 can be given analytical structure. Consider the annuitization decision at $\alpha = 0$. Shifting one dollar from liquid wealth to the annuity premium changes welfare by

$$\left. \frac{dV}{d\alpha} \right|_{\alpha=0} = W_0 \left[\text{MWR} \cdot V_A(W_0, 0, H, 1) - V_W(W_0, 0, H, 1) \right], \quad (15)$$

where V_W and V_A are the partial derivatives of the value function with respect to liquid wealth and annuity income. Define the *annuity value ratio*

$$R \equiv \text{MWR} \cdot \frac{V_A}{V_W}. \quad (16)$$

The agent annuitizes a positive fraction if and only if $R > 1$.

Proposition 1 (Yaari). *With no bequests ($\theta = 0$), deterministic health, fair annuity pricing, and $\beta(1+r) = 1$, the envelope theorem implies $V_W = u'(c^*)$ at the optimum. The annuity-income derivative V_A is the present discounted value of future marginal utilities generated by the income stream, weighted by survival probabilities. Consumption smoothing across states equates these marginal values, yielding $R = 1$ with full annuitization weakly optimal.*

Each demand-suppressing channel introduces a factor $\Phi_i < 1$ that reduces R below the Yaari benchmark:

Proposition 2 (Super-additive decomposition). *The annuity value ratio can be written*

$$R = \Phi_{\text{MWR}} \times \Phi_{\text{bequest}} \times \Phi_{\text{health}} \times \Phi_{\text{inflation}} \times \Phi_{\text{SS}} \times \Phi_{\text{pessimism}} \times \Phi_{\text{age}} \times \Phi_{\text{state}}, \quad (17)$$

where:

- $\Phi_{\text{MWR}} = \text{MWR}$. The pricing load directly scales the numerator. At $\text{MWR} = 0.87$, each dollar of premium buys 0.87 of expected value, a load of 13% relative to the actuarially

fair benchmark.

- $\Phi_{bequest} = \left[1 + (1-s) v'(b)/u'(c)\right]^{-1}$, where s is the survival probability and $v'(b)/u'(c)$ is the ratio of marginal bequest to marginal consumption utility. Under the DFJ luxury-good specification ($\kappa = \$272,628 \gg \text{median wealth} \approx \$40,000$), $v'(b) \approx \theta(\kappa)^{-\gamma}$ is nearly constant and small, so $\Phi_{bequest} \approx 1.00$.
- Φ_{health} reflects the covariance between survival and the marginal utility of liquid wealth. When health deterioration simultaneously reduces survival and raises medical costs, the ratio V_A/V_W falls. This is the Reichling–Smetters mechanism.
- $\Phi_{inflation}$ equals the ratio of the present value of the nominal payment stream (in real terms) to the real payment stream, weighted by marginal utilities. At 2% inflation, payments in year 25 retain only 61% of their initial purchasing power.
- Φ_{SS} captures the diminishing marginal insurance value of additional annuitization when Social Security already provides a longevity-insured income floor.
- $\Phi_{pessimism} < 1$ because the agent uses $\psi \cdot s(t, H)$ in the Bellman equation. The subjective present value of future annuity income is lower than the objective present value used by the insurer to set the price.
- $\Phi_{age} < 1$ because declining consumption needs $w(t)$ reduce the marginal utility of late-life income, which is precisely the income the annuity provides.
- $\Phi_{state} \leq 1$ because poor-health states generate lower marginal utility of consumption, and the annuity pays into those states.

The bundled behavioral wedge does not enter equation (17) as a smooth multiplicative factor on R . The wedge is applied externally as a multiplicative reduction on predicted ownership (equation 12), not as a value-function adjustment. The marginal-margin framework above therefore characterizes the model’s no-behavioral baseline; the wedge then reduces

predicted ownership proportionally without altering the agent’s participation-threshold computation. The Shapley decomposition in Section 4.4, which is computed on model-internal channels only, attributes the channels in the no-behavioral baseline.

Proposition 3 (Super-additivity of participation effects). *The fraction of agents annuitizing is $\Pr(\alpha^* > 0) = \Pr(R_i > 1)$, where R_i varies across agents due to heterogeneity in wealth, health, and Social Security income. In log space,*

$$\ln R_i = \sum_j \ln \Phi_j + \varepsilon_i, \quad (18)$$

where ε_i captures agent-level heterogeneity. Each additional channel shifts the distribution of $\ln R_i$ leftward by $|\ln \Phi_j|$. Agents near the participation threshold ($R_i \approx 1$, i.e., $\ln R_i \approx 0$) are disproportionately eliminated, so the ownership drop from adding a channel is larger when other channels have already thinned the right tail of the R_i distribution. This super-additive structure motivates the exact Shapley decomposition in Section 4.4, which attributes the total demand reduction to individual channels without dependence on the order in which channels are added.

The dominant term is $\Phi_{\text{MWR}} = 0.87$, which alone reduces the annuity value ratio by 13%. The near-unity Φ_{bequest} under the DFJ specification explains why bequests contribute only 0.9 pp in the Shapley decomposition: the luxury-good curvature parameter makes bequest utility irrelevant for median-wealth retirees.

2.13 Solution method

The model is solved by backward induction over the state space (W, A, H, t) . The wealth grid uses 80 points with power-function spacing (denser at low wealth) over $[0, \$3 \text{ million}]$, covering the 99.5th percentile of the HRS wealth distribution. The annuity income grid uses 30 points with cubic spacing to resolve small purchases, and health is discrete with three states. Consumption is optimized at each grid point using Brent’s method. Value

function interpolation across the wealth dimension uses piecewise linear interpolation with flat extrapolation.

The annuitization fraction α is optimized over a 101-point grid at age 65. Grid convergence is verified at the 9-node production quadrature rule: predicted ownership ranges from 25.6% at the medium grid (60×20) to 23.6% at the fine grid (100×40), bracketing the production-grid no-behavioral baseline; the 2 pp variation across the production-and-finer range is small relative to the UK retention sensitivity of the wedge factor ρ . The coarse grid (40×15) deviates more substantially and is reported only as a diagnostic. See Appendix Table 9.

3 Calibration and Validation

Table 1 summarizes the parameter values. Each is drawn from a published source or calibrated to match moments from the HRS.

Preferences. I set $\gamma = 2.5$, within the range of 1–3 estimated by Chetty (2006) from labor supply elasticities. Lockwood (2012) uses $\gamma = 2$; the slightly higher value reflects the richer model environment. Section 5 reports results for $\gamma \in \{1.5, 2.0, \dots, 5.0\}$.

Age-varying consumption needs. The decline rate $\delta_c = 0.02$ matches the central estimate from Aguiar and Hurst (2013), who found that nondurable consumption expenditure declines at roughly 2% per year in retirement after controlling for changes in household composition. This rate is applied uniformly across health states.

State-dependent utility. The health-utility weights $\varphi = [1.00, 0.92, 0.82]$ follow the mapping used by Reichling and Smetters (2015), based on the estimates of Finkelstein et al. (2013) from the HRS. Good-health individuals receive full weight; Fair and Poor health reduce the marginal utility of consumption proportionally.

Table 1: Model Parameters

Parameter	Symbol	Value	Source
<i>Preferences</i>			
Risk aversion	γ	2.5	Within Chetty (2006) range (1–3)
Discount factor	β	0.97	Standard
Bequest intensity	θ	56.96	Lockwood (2012)
Bequest shifter	κ	\$272,628	Lockwood (2012) ; De Nardi (2004)
<i>Demographics</i>			
Entry age		65	Retirement
Maximum age		110	Lockwood (2012)
Risk-free rate	r	0.02	Real
<i>Health and mortality</i>			
Health states		3	Good, Fair, Poor
Hazard multipliers	$\mu(H, t)$	See text	HRS mortality data (see text)
Quadrature nodes		9	Gauss–Hermite
Survival pessimism	ψ	0.960	O’Dea and Sturrock (2023)
<i>Medical expenditures</i>			
Log mean at 65 (Fair)	μ_m	7.037	Jones et al. (2018)
Annual growth		0.065	Jones et al. (2018)
Log std dev	σ_m	1.4	Jones et al. (2018)
Consumption floor	$\underline{c}$	\$6,180	SSI + state supplement
<i>Annuity market</i>			
Money’s worth ratio	MWR	0.87	Mitchell et al. (1999)
Fixed cost	F	\$2,500	Lockwood (2012)
Inflation rate	π	0.02	Post-Volcker average
<i>Preference and structural channels</i>			
Age-varying needs decline	δ_c	0.02	Aguiar and Hurst (2013)
Health-utility weights	$\varphi(H)$	[1.00, 0.92, 0.82]	Finkelstein et al. (2013)
Public-care aversion	χ_{LTC}	0.7	Ameriks et al. (2020)
<i>Bundled behavioral wedge (multiplicative transport)</i>			
UK retention factor (mid)	ρ	0.179	UK 2015 reform (Association of British Insurers, 2020); (Financial Conduct Authority, 2025)

All dollar values in 2014 dollars.

Public-care aversion (χ_{LTC}). The structural channel $\chi_{\text{LTC}} = 0.7$ is calibrated within the upper bound of the Ameriks et al. (2020 ECMA) 95% CI for public-care utility weight. Their central estimate is 0.5; we adopt 0.7 as a conservative choice that avoids overcalibrating this channel relative to other suppression mechanisms in the model. The channel applies a utility multiplier of $\chi_{\text{LTC}} < 1$ when the consumption floor binds and the agent is in Poor health (the Medicaid-LTC binding state), capturing retirees’ preference for self-financed over publicly-financed long-term care.

Bundled behavioral wedge (multiplicative transport). The bundled wedge—containing Force A (SDU), Force B (narrow-framing PED), and Force C (choice-architecture salience) as named conceptual components—is identified externally from the UK 2015 pension-freedoms reform and applied to the model’s no-behavioral baseline as a deterministic multiplicative transformation (Section 2.10, equation 12). The proportional retention factor at the production midpoint is $\rho = 0.179$ (UK post-reform retention 17% divided by pre-reform compulsory retention 95%). The UK retention sensitivity range of [13%, 25%] maps to a US prediction bracket of [5.7%, 11.0%] when applied to the no-behavioral baseline of 1.7%. The wedge is NOT a model parameter: the SDU consumption transformation and the at-purchase penalty mechanism have been removed from the model under Option 1 bundled identification. Forces A, B, and C are conceptually named but not separately parameterized; they are observed jointly in UK 2015 and transported jointly. The English Longitudinal Study of Ageing (ELSA) pension grid provides supporting individual-level evidence: in wave 6 (2012–13, pre-freedoms), 90.2% of DC pension recipients reported annuity-style income under the mandate; pooling waves 8–11 (2016–24, post-freedoms), 1.27% of observed lump-sum disposition records used the lump sum to buy annuity income (n=869).

Out-of-sample reading. The calibration is set entirely from non-US data. Both HRS measures of US lifetime annuity ownership—the conventional any-annuity income proxy (3.34%, 95% CI [2.89%, 3.85%]) and the cleaner fat-file lifetime-contract indicator (2.02%, 95% CI

[1.68%, 2.43%])—fall within the wider sensitivity range. The lifetime contract indicator is just below the lower edge of the headline retention bracket [4.6%, 8.8%]; the income proxy is at the lower edge. No US moment is used to set the calibration; consistency with US ownership is an out-of-sample comparison.

The bequest parameters $\theta = 56.96$ and $\kappa = \$272,628$ are taken directly from [Lockwood \(2012\)](#), who estimated them from HRS bequest data at $\gamma = 2$. Using them at $\gamma = 2.5$ raises a portability concern in principle, but simulation-based checks show less than 14% divergence in the bequest-to-wealth ratio across $\gamma \in [1.5, 5.0]$, below the threshold warranting recalibration. All results use Lockwood’s original estimates.

Health and mortality. Health transition matrices are calibrated to HRS panel data following [De Nardi et al. \(2010\)](#). The five HRS self-reported health categories are mapped to three states: Good (Excellent/Very Good), Fair (Good), and Poor (Fair/Poor). Baseline survival probabilities use the SSA administrative life table from [Lockwood \(2012\)](#).

Hazard multipliers are estimated from the RAND HRS longitudinal file (waves 4–16, $N = 126,249$ person-wave observations aged 65+). The estimated ratios by age band are:

Age band	Good	Fair	Poor
65–74	0.49	1.00	3.29
75–84	0.60	1.00	2.77
85+	0.74	1.00	1.82

The compression of the gradient at older ages is consistent with selection effects: those surviving to 85 in Poor health have already demonstrated above-average resilience. The baseline constant multipliers [0.50, 1.0, 3.8] represent a central tendency across these age bands, falling between the R-S functional estimates ([0.45, 1.0, 3.5], consistent with ADL/IADL-based measures) and the empirical HRS self-reported health values ([0.57, 1.0, 2.7]). [Section 5](#) reports results across the full range.

Medical expenditures. The medical expenditure process matches the moments published by Jones et al. (2018): mean out-of-pocket spending of \$4,200 at age 70, growing to \$29,700 at age 100, with a 95th percentile of approximately \$111,200 at age 100. Health shifts both the mean and variance of log spending, with Poor health raising expected costs by a factor of roughly two relative to Fair.

Population sample and empirical targets. The wealth distribution is drawn from the RAND HRS longitudinal file, selecting single retirees aged 65–69 using filters comparable to Lockwood (2012). The resulting sample contains 5,303 person-wave observations (waves 5–9). The main analysis restricts to the 2846 observations with liquid wealth above \$5,000, as those with negligible wealth cannot meaningfully annuitize. *The empirical-target ownership rates reported below (2.02% lifetime contract indicator and 3.34% any-annuity income proxy) are computed on the same wealth-restricted subsample as the model predictions, ensuring symmetric comparison.* The structural model maps each household’s state to a predicted decision rather than estimating moments from the sample distribution. The HRS observations are unweighted; the qualitative results are stable under sample weights. The any-annuity income-proxy ownership rate is 3.36% unweighted and 3.70% under RAND HRS sample weights—a 0.35 pp gap that leaves the model’s bracket overlap unchanged in both directions. Unweighted estimates serve as the headline so that the empirical target maps directly to the cells solved in the structural decomposition.

The paper reports two HRS measures of US lifetime annuity ownership in parallel as empirical targets:

1. *Lifetime annuity contract indicator (q286 series; preferred measure).* Drawn from the HRS pension grid (“fat-file”) waves 5–9. Identifies respondents reporting at least one annuity contract for which the question stem “Will this annuity continue for the rest of your life?” is answered affirmatively. In the pooled sample of 5,396 eligible person-waves, 109 report a lifetime annuity contract: 2.02% (Wilson 95% CI [1.68%, 2.43%]).

This is the methodologically appropriate target for an SPIA-focused model.

2. *Any-annuity income proxy ($r\{w\}iann$; conventional measure).* The RAND HRS variable $r\{w\}iann$ flags any positive annuity income in the wave, including DC pension withdrawals and short-period payouts not life-contingent. Pooled across waves 5–9, 3.34% of person-waves report positive annuity income (Wilson 95% CI [2.89%, 3.85%]). This measure—or a closely comparable any-annuity-income construct—is used by most prior literature (e.g., Lockwood, 2012 reports a 3.6% rate for single retirees 65–69 using a wave-specific HRS annuity-income measure that conceptually overlaps with $r\{w\}iann$ but is not identical at the variable-naming level).⁴ It overstates lifetime annuity ownership by conflating it with one-time DC pension withdrawals.

Both HRS measures fall within the model’s UK-anchored sensitivity range of approximately [2.5%, 8.8%], with the conventional income proxy near the lower-middle and the cleaner lifetime-contract indicator just below the lower edge of the headline UK retention bracket of [4.6%, 8.8%]. Section 4 reports headline results against both measures.

Lifecycle moment validation. Table 2 reports a compact bequest-distribution diagnostic from the lifecycle simulation: simulated vs. HRS exit-interview moments for the mean, median, and fraction with bequest greater than \$10K. These moments are not targeted in calibration; the comparison serves as external discipline on the bequest and decumulation behavior rather than as matched moments.

The model understates bequests in this representative simulation: the simulated median is zero versus an HRS exit-interview target of \$20,000, and the simulated mean (\$28,657) is well below the empirical mean (\$90,000). The simulated fraction with bequest above \$10K (26.1%) is below the HRS rate (45.0%). This reinforces that the welfare and ownership

⁴Lockwood’s measure is constructed from wave-specific HRS pension and income variables; the RAND HRS Longitudinal File’s $r\{w\}iann$ is the canonical pooled variable that subsequent literature has used as the analogous measure across waves. The two yield essentially identical pooled rates for single retirees 65–69, but readers seeking exact variable mapping to Lockwood’s construct should consult his data appendix.

results should not be interpreted as a full bequest-distribution fit. The DFJ luxury-good calibration concentrates bequest motives at high wealth, so the median household in the representative simulation (initial wealth \$250,000, Fair health) does not retain wealth for bequests after correlated medical and longevity risk. A full bequest-targeted calibration would change θ and κ jointly; the [Lockwood \(2012\)](#) estimates used here are kept fixed in the spirit of the unified-channel decomposition exercise.

Table 2: Simulated vs Empirical Lifecycle Moments

Moment	Simulated	Empirical (HRS)
Mean bequest	\$40023	\$90000
Median bequest	\$0	\$20000
Fraction bequest > \$10K	37.0%	45.0%

Simulated: 100,000 trajectories, initial wealth \$250,000, Fair health.

Empirical: HRS exit interviews (bequests); Jones et al. (2018) (medical).

4 Results

4.1 Sequential decomposition

Table 3 presents an intuitive sequential decomposition of predicted voluntary annuity ownership, adding channels one at a time across rational, preference, and behavioral blocks. Because this exercise is order-dependent, I use it as a descriptive path through the model and treat the exact Shapley values below as the preferred attribution. For narrative transparency the table reports medical expenditure risk and the health-mortality correlation as two separate sequential steps; the Shapley analysis combines them into a single Med+R-S channel (Section 4.4), because the R-S mechanism’s quantitative bite in this framework operates through the interaction with medical risk: without competing demand for liquid wealth in sick states, the lower expected annuity NPV does not translate into a precautionary motive against annuitization. The narrative below describes the rational channels in this section and the preference/behavioral channels in Sections 4.2 and 4.3; the table covers all

three blocks for ease of reference.

Table 3: Sequential Decomposition of Predicted Voluntary Annuity Ownership

Channel	Ownership (%)	Δ (pp)	Retention	Cumulative
Yaari benchmark	41.8	—	—	—
+ Social Security	100.0	+58.2	239.0%	2.3896
+ Bequest motives	100.0	+0.0	100.0%	2.3896
+ Medical expenditure risk (uncorrelated)	97.4	-2.6	97.4%	2.3266
+ Health-mortality correlation (R-S)	87.7	-9.6	90.1%	2.0966
+ Survival pessimism	62.4	-25.3	71.1%	1.4912
+ State-dependent utility	62.7	+0.3	100.5%	1.4979
+ Age-varying consumption needs	54.9	-7.8	87.6%	1.3123
+ Realistic pricing loads	8.3	-46.6	15.1%	0.1982
+ Inflation erosion	1.8	-6.5	21.2%	0.0420
Observed (Lockwood 2012)	3.6	—	—	—

Retention rate = ownership after channel / ownership before channel. Cumulative product of retention rates tracks geometric compounding.

Population benchmark with public consumption floor. The first row—no bequest motive, actuarially fair pricing, no medical expenditure risk, no inflation, and no Social Security income—predicts 41.8% ownership in the HRS population sample. This is not the theoretical Yaari full-annuitization result. The benchmark holds two US institutional features fixed across all subsequent decomposition rows: a public consumption floor (Medicaid/SSI proxy at $c_{\text{floor}} = \$6,180/\text{year}$) and a minimum-wealth analytical sample restriction (HRS singles with non-housing wealth $\geq \$5,000$). Many low-wealth households in the data find annuitization unattractive even in this frictionless environment because the public floor effectively guarantees consumption at Medicaid-eligible levels, reducing the marginal value of liquid wealth that could be annuitized. The standard convention in the calibrated lifecycle literature (Pashchenko 2013; Peijnenburg-Nijman-Werker 2016; De Nardi-French-Jones 2010) is to bake the safety net into the institutional environment rather than treat it as a behavioral channel; this paper follows that convention.

Social Security. Adding Social Security income (COLA-protected, calibrated by wealth quartile) raises predicted ownership to 100.0%. SS acts as a complement in the frictionless model: the guaranteed income floor enables consumption smoothing and makes agents willing to lock up additional wealth in annuities. The 239.0% retention rate reflects this complementarity.

Bequest motives. Adding the DFJ luxury-good bequest specification ($\theta = 56.96$, $\kappa = \$272,628$) leaves ownership unchanged at 100.0% (retention rate of 100%). The luxury-good $\kappa = \$272,628$ makes marginal bequest utility nearly flat for individuals with wealth below \$200,000. In the HRS sample, median liquid wealth is approximately \$40,000.

Medical expenditure risk and health-mortality correlation (combined). Introducing stochastic medical expenditures together with the Reichling–Smetters health-mortality correlation reduces ownership to 87.7% (retention 87.7%, a -12.3 pp effect). The two are bundled into a single channel because the R-S mechanism’s quantitative bite in this framework operates through the interaction with medical risk: without competing demand for liquid wealth in sick states, the lower expected annuity NPV does not translate into a precautionary motive against annuitization. The mechanism: when a retiree’s health deteriorates, expected remaining lifetime falls (reducing the present value of future annuity payments) and expected medical costs rise (increasing the marginal value of liquid wealth). The demand-suppressing effect depends on this correlation between health deterioration, survival reduction, and cost increases. The combined channel was not jointly incorporated in the multi-channel models of [Pashchenko \(2013\)](#) or [Peijnenburg et al. \(2016\)](#).

Survival pessimism. Introducing subjective survival beliefs ($\psi = 0.960$) reduces ownership from 87.7% to 62.4% (retention 71.1%, a -25.3 pp effect). The agent overweights the probability of dying before recovering the premium.

Pricing loads. Repricing the annuity at $MWR = 0.87$ and adding a \$2,500 fixed purchase cost produces the largest single reduction, from 62.4% to 18.2% (retention 29.1%). An 13% load eliminates the surplus that marginally willing buyers retained after accounting for health risk and survival pessimism.

Inflation erosion. Converting the annuity from real to nominal at 2% annual inflation reduces ownership from 18.2% to 8.5% (retention 46.6%). This step is a product counterfactual: the previous steps priced a real annuity (constant purchasing power); this step switches to a nominal annuity whose real value erodes over time. Social Security income is COLA-protected and does not erode, but the private annuity payment loses purchasing power: at 2% inflation, real income from the annuity falls by 39% over 25 years.

Summary. The six standard rational channels predict 8.5% ownership relative to a frictionless population benchmark of 41.8%. Pricing loads (retention 29.1%) and the combined Med+R-S channel (retention 87.7%, reported as two separate steps in Table 3 for narrative transparency) are the two most powerful channels in this sequential ordering. The key implication is that the standard rational literature substantially overshoots both empirical targets—a residual gap that the preference and behavioral channels in subsequent sections address.

4.2 Preference Channel Extensions

Table 4 isolates the incremental effect of the two added preference channels.

Adding age-varying consumption needs ($\delta_c = 0.02$, calibrated to [Aguiar and Hurst, 2013](#)) reduces predicted ownership from 8.5% to 2.6%. This is the quantitatively important extension beyond the six-channel rational model: if late-life consumption needs decline, the insurance value of a level annuity stream declines as well. The resulting reduction is comparable in magnitude to survival pessimism or inflation erosion.

Table 4: Two-Model Decomposition: Structural Channels and UK Reduced-Form Transport

Specification	Ownership (%)	Δ (pp)
Six rational channels (Layer 1)	8.5	—
+ Age-varying consumption needs	2.6	-5.8
+ State-dependent utility	1.8	-0.9
+ Public-care aversion χ_{LTC} (Layer 2 complete)	1.7	-0.1
+ Source-dependent utility (Force A)	27.8	+26.1
+ Narrow-framing penalty (Force B; Model 1)	0.0	-27.8
Model 2: frictionless baseline (41.8) $\times$ UK 17/95	7.5	—

Model 1 (top block) is the structural multi-channel decomposition. Layer 1 covers rational frictions; Layer 2 adds preference and structural channels; Force A (SDU) and Force B (PED) close the model with behavioral channels calibrated to Blanchett-Finke (2024-25) and Chalmers-Reuter (2012) respectively.

Model 2 (bottom row) applies the UK 2015 pension-freedoms retention factor (UK post / UK pre = 0.18 at the production midpoint) to the frictionless Yaari baseline. The UK retention range [13%, 25%] maps to a Model 2 prediction bracket of [5.7%, 11.0%].

Adding state-dependent utility ($\varphi = [1.00, 0.92, 0.82]$, taken directly from the central estimates of [Finkelstein et al., 2013](#)) produces only a small additional reduction, from 2.6% to 1.5%. The channel’s effect is insensitive to the specific mapping within the range the literature supports: under the softer translation $\varphi = [1.00, 0.95, 0.85]$ used by [Reichling and Smetters \(2015\)](#), the full-model prediction shifts to 2.0%, a difference of 0.5 pp. In either calibration, this channel is best interpreted as a completeness check rather than a load-bearing mechanism.

The full eight-channel rational+preference model therefore predicts 1.5% ownership under the raw FLN central mapping and 2.0% under the Reichling–Smetters translation, a reduction from the 41.8% frictionless benchmark but still substantially above both empirical targets (3.34% on the conventional any-annuity income proxy; 2.02% on the cleaner life-time contract indicator). The remainder of the paper uses the raw FLN mapping as the production calibration.

4.3 Behavioral channel extensions: the two-force structure

The eight-channel model (six rational plus two preference channels) predicts 1.8% ownership, substantially above both empirical targets (2.02% on the lifetime contract indicator, 3.34% on the conventional any-annuity income proxy). The standard rational and preference account, on its own, does not dissolve the puzzle. The paper’s central structural addition is the introduction of two behavioral channels operating at distinct decision moments.

The bundled behavioral wedge. The bundled wedge is applied to the model’s no-behavioral baseline as a deterministic multiplicative transformation (Section 2.10, equation 12): predicted US ownership equals the no-behavioral baseline multiplied by the UK proportional retention factor $\rho = \text{post/pre}$. At the production midpoint ($\rho = 17/95 \approx 0.18$), applying the wedge to the no-behavioral baseline of 1.7% yields a production headline prediction of 7.5%. The wedge is post-processing on ownership rates; no behavioral parameter is calibrated inside the Bellman equation.

Sensitivity bracket from UK retention range. Uncertainty in the UK post-reform retention rate (13–25% across ABI/FCA estimates) maps directly to a US prediction bracket of [5.7%, 11.0%]. A wider sensitivity range across alternative anchors, including the descriptive ELSA microdata total drop, spans approximately [2.5%, 8.8%]. The bracket reflects uncertainty in the UK identifying moment, not parameter calibration uncertainty inside the model.

Bracket against the empirical targets. The conventional any-annuity income proxy (3.34%, Wilson 95% CI [2.89%, 3.85%]) and the cleaner lifetime-contract indicator (2.02%, Wilson 95% CI [1.68%, 2.43%]) are reported as out-of-sample comparisons against the wedge bracket; neither bracket bound is tuned to a US moment.

Decomposition within the bundle is not data-identified. The bracket reflects the joint footprint of Forces A, B, and C; the data permit transporting only the bundle, not the individual channel contributions. Counterfactuals that “remove a single force” are therefore reported in Section 4.10 as illustrative bundle-removal exercises, not as separately-identified channel attributions. Under joint identification (Section 2.10), sub-decomposing the wedge into Force A / Force B / Force C contributions is parameter-dependent rather than data-identified.

4.4 Exact Shapley decomposition

The sequential decomposition is order-dependent: the contribution attributed to each channel depends on which channels precede it. To obtain order-independent attributions, I compute exact Shapley values by solving the model for all $2^{10} = 1,024$ channel subsets and averaging each channel’s marginal contribution across all possible orderings. Medical expenditure risk and the Reichling–Smetters health-mortality correlation are bundled into a single channel, because the R-S mechanism’s quantitative bite in this framework operates through the interaction with medical risk: without competing demand for liquid wealth in sick states, the lower expected annuity NPV does not translate into a precautionary motive against annuitization; this reformulation also removes the non-standard coupling rule that the prior 11-channel structure required.

Table 5 reports the exact order-independent attribution results for the nine model-internal channels (the bundled behavioral wedge is applied separately as multiplicative post-processing and does not appear in the Shapley enumeration). Among demand-suppressing channels, pricing loads (14.0 pp, 33%) and the combined Med+R-S health-mortality channel (8.6 pp, 21%) are the two largest contributors, followed by age-varying consumption needs (2.9 pp, 7%), survival pessimism (5.9 pp, 14%), and inflation erosion (1.4 pp, 3%). Bequest motives (0.9 pp) and state-dependent utility (0.3 pp) are quantitatively negligible.

Two channels enter with negative Shapley values, reflecting their role as ownership boost-

ers: source-dependent utility (Force A) at -13.4 pp (-32%) and Social Security at -17.6 pp (-42%). Source-dependent utility raises annuitization by transforming portfolio wealth into the income mental account at full utility weight; Social Security raises annuitization by providing the income floor that makes additional annuitization rational at the margin. The shares sum to more than 100% in absolute value because channels interact super-additively and because the negative-shapley boosters partly offset the suppressors.

The bundled behavioral wedge does not appear in the Shapley enumeration. Under Option 1 bundled identification, the wedge is applied as multiplicative post-processing on ownership rates (equation 12), not as a model channel: the Shapley enumeration covers only the model-internal nine channels listed above. The wedge’s contribution is therefore reported separately as the multiplicative reduction from 1.7% (the nine-channel no-behavioral baseline) to 7.5% (production midpoint after wedge), with bracket [5.7%, 11.0%].

Table 5: Exact Shapley-Value Decomposition of Predicted Annuity Ownership

Channel	Shapley (pp)	Share (%)
SS	-17.65	-42.2
Bequests	+0.85	2.0
Medical+R-S	+8.60	20.6
Pessimism	+5.85	14.0
Age needs	+2.90	6.9
State utility	+0.26	0.6
Loads	+14.01	33.5
Inflation	+1.44	3.4
Public-care aversion (LTC)	-0.64	-1.5
Source-dependent utility (SDU)	-13.44	-32.1
Narrow-framing penalty (PED)	+39.65	94.7
Total	+41.85	100.0

Exact Shapley values computed from all 2048 channel subsets. Each value represents the weighted average marginal ownership reduction (pp) across all coalition orderings. Yaari baseline: 41.8%. Full model: 0.0%.

The Shapley values sharpen three findings from the sequential decomposition. First, pricing loads and the combined Med+R-S channel are the two largest demand-suppressing contributors regardless of ordering. Second, the channels not jointly incorporated in the

earlier multi-channel literature—the combined Med+R-S correlation, survival pessimism, and age-varying needs—collectively account for 17.4 pp (41%), comparable to pricing loads. Third, bequests contribute non-trivially (0.9 pp) despite their zero sequential retention rate, because order-independent averaging captures the states in which bequests matter only after pricing and inflation are already active.

4.5 Pairwise channel interactions

Table 6 reports the pairwise interaction strength for the eight rational and preference channels. For each pair (A, B), the interaction is defined as the difference between the ownership drop when both are active simultaneously and the sum of their individual drops. A negative interaction indicates super-additivity: the channels reinforce each other. The bundled behavioral wedge is omitted from this table because it is applied as multiplicative post-processing on the no-behavioral baseline rather than as a model channel; the order-independent contributions of the model-internal channels are reported in the exact Shapley decomposition (Section 4.4).

Table 6: Pairwise Interaction Strengths, Rational and Preference Channels (pp)

	SS	Bequests	Medical+R-S	Loads	Inflation	Pessimism	HealthUtil	AgeNeeds
SS	—	+0.0	-1.3	-31.3	+2.3	-11.5	+0.2	-0.4
Bequests	—	—	-0.3	-0.1	-0.0	-0.1	+0.0	-0.1
Medical+R-S	—	—	—	-2.5	+0.8	-0.6	+0.4	-0.6
Loads	—	—	—	—	-0.1	-0.3	+0.0	+0.0
Inflation	—	—	—	—	—	+0.9	+0.0	+0.4
Pessimism	—	—	—	—	—	—	+0.2	+0.1
HealthUtil	—	—	—	—	—	—	—	-0.1
AgeNeeds	—	—	—	—	—	—	—	—

Each cell shows the interaction: ownership with both channels minus the additive prediction from individual effects. Negative values indicate super-additive demand reduction (channels reinforce each other).

The strongest interaction is between Social Security and pricing loads (−21.1 pp). SS raises demand by providing an income floor, and loads destroy demand by eliminating the

surplus; when combined, the effect of loads is amplified by the elevated demand that SS creates. The combined Med+R-S health-mortality channel and pricing loads interact at -2.4 pp, reflecting the economic complementarity between valuation risk and pricing frictions. Bequest interactions are near zero across the board, consistent with the modest isolated effect of the DFJ luxury-good specification.

4.6 Comparison with prior work

Table 7 evaluates the channels that [Pashchenko \(2013\)](#) emphasized—bequest motives, minimum purchase requirements, and pricing loads—within the present framework. This is not a replication of her model, which differs in preference specification, health dynamics, and solution method. Rather, it quantifies how much of the ownership gap those channels account for in the unified model, and how much additional reduction the present paper obtains from the broader channel set. Her model predicted approximately 20% participation; the six-channel rational model here predicts 8.5%, and the added preference channels bring the prediction to 1.8%—still substantially above both empirical targets, with the residual closed only by the two behavioral channels (Force A + Force B).

[Peijnenburg et al. \(2016\)](#) found that full annuitization remains approximately optimal in their framework. Their model includes background risk and default risk but omits the health-mortality correlation that drives the Reichling–Smetters mechanism. Without the specific correlation between health shocks, survival, and medical costs, medical expenditure risk alone may increase annuity demand. The combined Med+R-S channel in the present model contributes 12.3 pp in the sequential decomposition (and 8.6 pp in the order-independent Shapley attribution), confirming that the correlation is a quantitatively important channel that prior work omitted.

Table 7: Pashchenko (2013) Channels in Our Framework

Model Specification	Ownership (%)	Δ
Yaari benchmark (SS on)	100.0	—
+ Bequest motives (DFJ)	100.0	+0.0 pp
+ Minimum purchase (25K)	74.1	-25.9 pp
+ Pricing loads (MWR=0.85)	59.5	-14.7 pp
<i>Channels omitted by Pashchenko:</i>		
+ Medical costs + R-S correlation	38.2	-21.3 pp
+ Survival pessimism	13.7	-24.4 pp
+ Full loads + Inflation	8.5	-5.3 pp
Observed (Lockwood 2012)	3.6	—

Steps 0–3 incorporate the channels Pashchenko (2013) identified (SS, bequests, minimum purchase, pricing loads) into our unified model. Steps 4–6 add channels not in her framework. Note: this is not a replication of her model (different preferences, health dynamics, solution method). Housing illiquidity is not modeled; see text.

4.7 Sensitivity to risk aversion

The decomposition results depend on calibrated parameters, particularly risk aversion γ . Table 8 reports predicted ownership across a range of γ values and inflation rates for the six-channel rational model.

Table 8: Predicted Ownership (%) by Risk Aversion and Inflation Rate

	$\pi = 1\%$	$\pi = 2\%$	$\pi = 3\%$
$\gamma = 2.4$	0.3	0.1	0.0
$\gamma = 2.5$	4.7	1.8	0.7
$\gamma = 2.6$	9.6	6.4	3.3
$\gamma = 3.0$	22.5	18.1	13.6

Baseline: $\gamma = 2.5$, $\pi = 2\%$, DFJ bequests, MWR = 0.87, hazard multipliers [0.50, 1.0, 3.0].

At the baseline $\gamma = 2.5$, the six-channel rational model predicts 8.5%. Predicted ownership rises with γ : 0.0% at $\gamma = 2.0$, 0.1% at $\gamma = 2.4$, and 18.1% at $\gamma = 3.0$. The eight-channel rational+preference model (with age-varying needs and state-dependent utility) predicts 1.8% at the baseline $\gamma = 2.5$, still substantially above both empirical targets. Adding the

structural public-care aversion channel (χ_{LTC}) yields the no-behavioral baseline of 1.7%; applying the bundled behavioral wedge as multiplicative transport (Section 2.10) yields the production headline prediction of 7.5% at the UK production midpoint with bracket [5.7%, 11.0%] across UK retention anchors.

The sensitivity to γ is not unique to this model. It is shared by all lifecycle annuitization models and reflects the near-margin-of-indifference nature of the annuity purchase decision. The mortality credit is modest for healthy 65-year-olds, so the net surplus from annuitization sits near zero for many households. Small changes in risk aversion tip the optimal decision from non-purchase to purchase.

A multi-gamma diagnostic decomposition (run at $\gamma = 2.0, 2.5$, and 3.0) confirms that the first five channels (Yaari through pessimism) are gamma-insensitive: all three decompositions reach 71–80% ownership after the pessimism step. The sensitivity concentrates in pricing loads and inflation erosion. Higher risk aversion generates stronger insurance demand, which survives pricing frictions for more agents.

4.8 Joint parameter uncertainty

The single-parameter sweeps above hold all other parameters at their calibrated values. To assess robustness to joint calibration uncertainty within the model-internal parameters, I draw 1,000 parameter vectors from independent uniform distributions on the empirically supported ranges: $\mu_P \sim U(2.0, 3.5)$, $\pi \sim U(0.015, 0.025)$, $\text{MWR} \sim U(0.83, 0.91)$, $\psi \sim U(0.97, 1.0)$, $\delta_c \sim U(0.01, 0.03)$. Risk aversion is held fixed at the baseline. For each draw the model is solved on the production grid and the no-behavioral ownership prediction recorded; the bundled behavioral wedge is then applied as multiplicative transport.

The no-behavioral baseline prediction sits within a 90% sensitivity interval of [17.9%, 43.3%] (median 31.9%, IQR [26.5%, 36.5%]) under joint parameter uncertainty. The interval is reasonably tight because the joint draw is over five rational/preference parameters whose effects partly offset each other. The bundled-wedge prediction inherits this uncertainty

Table 9: Conditional Monte Carlo: Predicted Ownership at $\gamma = 2.5$

Statistic	Value
Number of draws	1000
Median predicted ownership	31.9%
90% sensitivity interval	[17.8%, 43.3%]
Interquartile range (50% interval)	[26.5%, 36.5%]
Mean	31.4%
Min / Max	7.0% / 53.2%
Fraction in [1%, 10%]	0%
Fraction in [3%, 6%] (observed range)	0%

Risk aversion fixed at $\gamma = 2.5$. Joint draws over six calibration-uncertain parameters: $\mu_P \sim U(2.0, 3.5)$, $\pi \sim U(0.015, 0.025)$, $MWR \sim U(0.83, 0.91)$, $\psi \sim U(0.97, 1.0)$, $\delta_c \sim U(0.01, 0.03)$, $\psi_{\text{purchase}} \sim U(0.005, 0.030)$. Full ten-channel model.

multiplicatively: applying the production midpoint factor $\rho = 0.179$ to the 90% interval yields a wedged-prediction interval that compounds the model’s parameter uncertainty with the UK retention sensitivity range reported in Section 4.3. Both HRS empirical targets (2.02% on the lifetime contract indicator, 3.34% on the conventional any-annuity income proxy) are reported as out-of-sample comparisons against the wedged-prediction range.

4.9 Heterogeneous welfare

Table 10 reports the consumption-equivalent variation (CEV) from annuity market access under baseline pricing, across household types defined by initial wealth, health status, and bequest motive intensity.

Under the conservative (ABI-anchored) bracket calibration with all ten channels active, the welfare gain from annuity market access is small at every wealth level. The largest single-cell CEV is 0.0% at \$100,000 in Good health under no bequests; with the DFJ bequest specification it is 0.0% at the same wealth and health. CEV at \$200,000 in Good health is 0.0% (no bequest) and 0.0% (DFJ). At wealth above \$500,000 and at all health levels Fair or Poor, CEV is essentially zero. At the population level, mean CEV under DFJ bequests is 0.0%, with 9.2% of households deriving any positive welfare gain and 0.0% deriving gains

Table 10: Consumption-Equivalent Variation by Household Type

Wealth	Health	No bequest	Moderate (DFJ)	Strong bequest
\$10,000	Good	0.00%	0.00%	0.00%
	Fair	0.00%	0.00%	0.00%
	Poor	0.00%	0.00%	0.00%
\$25,000	Good	0.00%	0.00%	0.00%
	Fair	0.00%	0.00%	0.00%
	Poor	0.00%	0.00%	0.00%
\$50,000	Good	0.00%	0.00%	0.00%
	Fair	0.00%	0.00%	0.00%
	Poor	0.00%	0.00%	0.00%
\$100,000	Good	0.00%	0.00%	0.00%
	Fair	0.00%	0.00%	0.00%
	Poor	0.00%	0.00%	0.00%
\$200,000	Good	0.00%	0.00%	0.00%
	Fair	0.00%	0.00%	0.00%
	Poor	0.00%	0.00%	0.00%
\$500,000	Good	3.20%	0.24%	0.00%
	Fair	2.96%	0.07%	0.00%
	Poor	2.50%	0.00%	0.00%
\$1,000,000	Good	7.72%	0.86%	0.00%
	Fair	7.07%	0.59%	0.00%
	Poor	5.87%	0.14%	0.00%

CEV: consumption-equivalent variation (%). Positive values indicate welfare gain from annuity market access. Full model with medical costs, health-mortality correlation, MWR = 0.87, inflation = 2%, $\gamma = 2.5$.

exceeding 1% of consumption.

The economic content of these small CEVs is the same as the low predicted ownership rate: most households would not optimally purchase an annuity at current prices, so giving them access to the market provides little marginal value. The CEV grid is therefore the welfare counterpart of the ownership prediction—both confirm that the marginal household is approximately indifferent at the no-purchase margin. Larger welfare gains appear only when policy changes the optimal choice: lowering the load (group pricing) or removing the bundled behavioral wedge (default architecture), reported in Section 4.10.

Section 4.10 examines how supply-side reforms change these welfare results.

Normative interpretation of the behavioral wedge. The bundled wedge does not enter the within-model welfare computation directly: under Option 1 bundled identification, the wedge is applied as multiplicative post-processing on ownership rates rather than as a structural parameter affecting agent decisions. The agent’s CRRA value-ratio CEV is therefore computed on the no-behavioral baseline. A default-architecture intervention that “removes the wedge” has a normatively contingent welfare interpretation (Bernheim and Rangel, 2009; Beshears et al., 2008): under a preference interpretation of the underlying SDU/PED/salience phenomena (welfare-relevant), removing the wedge does not increase agent welfare because the agents’ own preferences include those phenomena; under a bias interpretation (welfare-irrelevant), removing the wedge captures the full rational-utility gain. The ranking of policy interventions across the bracket is preserved across both interpretations, but the magnitude of the demand-side default-architecture gain depends on which interpretation is adopted. The empirical exercise in this paper does not adjudicate the preference-vs-bias choice; welfare claims are reported with this caveat made explicit.

4.10 Policy counterfactuals

Table 11 reports predicted ownership under twelve policy scenarios that vary pricing, inflation protection, survival beliefs, and Social Security generosity. Table 13 reports the corresponding welfare gains.

Table 11: Predicted Annuity Ownership Under Policy Counterfactuals

Policy Scenario	MWR	Inflation	Ownership (%)	Δ (pp)
Baseline	0.87	2%	1.7	—
Group pricing (MWR=0.90)	0.90	2%	5.8	+4.2
Public option (MWR=0.95)	0.95	2%	17.6	+15.9
Actuarially fair (MWR=1.0)	1.00	2%	30.3	+28.6
Real annuity, TIPS-backed	0.78	0 (real)	0.0	-1.7
Real annuity, nominal-equiv	0.87	0 (real)	8.7	+7.1
Fair + real	1.00	0 (real)	37.1	+35.5
SS cut 23Correct pessimism ($\psi=1.0$)	0.87	2%	25.1	+23.4
Group + correct pessimism	0.90	2%	33.6	+31.9
Public consumption floor doubled	0.87	2%	0.7	-1.0
Best feasible package	0.90	0 (real)	42.6	+40.9
Default architecture (no wedge)	0.87	2%	1.7	+0.0
Default + group pricing	0.90	2%	5.8	+4.2
Observed (Lockwood 2012)			3.6	

Baseline: $\gamma = 2.5$, $\beta = 0.97$, DFJ bequests ($\theta = 56.96$, $\kappa = \$272,628$), $\psi = 0.981$.
Population: HRS single retirees 65–69 with $W \geq \$5,000$ ($N = 566$). Group pricing reflects TSP/employer plan MWR (James et al. 2006). SS cut: 23% across-the-board (projected trust fund exhaustion).

Pricing loads as a dominant supply-side lever. Group annuity pricing (MWR = 0.90), achievable through employer plans or government-sponsored pools (James and Song, 2006), raises predicted ownership from the conservative bracket end to 5.8%. A public option at MWR = 0.95 reaches 17.6%, and an actuarially fair benchmark at MWR = 1.00 reaches 30.2%. These results quantify the pricing channel’s importance: a modest improvement in the money’s worth ratio—from 0.87 to 0.90—raises participation by roughly nine percentage points relative to the conservative bracket end.

Inflation protection. A real annuity at the same MWR (0.87) leaves ownership essentially unchanged at 8.8%, while a TIPS-backed real annuity at $\text{MWR} = 0.78$ falls to 0.0%, illustrating that the pricing penalty for inflation indexing offsets most of the gain from removing inflation erosion. Combining inflation protection with fair pricing reaches 37.1%.

Social Security and the complement-substitute duality. A 23% across-the-board SS benefit cut—the projected consequence of trust fund exhaustion around 2033—raises private annuity demand from the conservative bracket end to 10.4%. Table 12 reports the full schedule of SS cuts. This is a stylized across-the-board reduction; actual benefit cuts would vary by claiming history and cohort. The relationship is non-monotonic: moderate cuts reduce crowd-out and increase private annuitization, but the largest cuts eventually lower ownership because they also remove the income floor that makes annuitization feasible for low-wealth households.

Table 12: Private Annuity Demand Response to Social Security Benefit Reductions

SS Benefit Cut	Ownership (%)	Mean α	Δ (pp)
0% (baseline)	1.8	0.002	—
10%	3.8	0.005	+2.1
15%	5.9	0.007	+4.1
23% (trust fund)	10.4	0.015	+8.7
30%	16.3	0.030	+14.5
40%	25.1	0.059	+23.3
50%	28.6	0.084	+26.8
100% (elimination)	14.3	0.067	+12.5
Observed (Lockwood 2012)	3.6		

Full model with all channels active. SS quartile levels scaled by $(1 - \text{cut fraction})$. Baseline levels: [\$14K, \$17K, \$20K, \$25K]. 23% cut corresponds to projected trust fund exhaustion circa 2033. 100% cut is a theoretical benchmark (complete SS elimination).

Best feasible (supply-side + information) package. For ownership counterfactuals, the “best feasible” scenario combines group pricing ($\text{MWR} = 0.90$), inflation protection (real annuity), and corrected survival beliefs ($\psi = 1.0$), producing 42.5% ownership at the

no-behavioral baseline. This is the upper bound of the supply-side and information interventions on the model-internal channels, not a realistic forecast of market participation. The “best feasible + default architecture” scenario in the welfare counterfactuals (Section 4.9) below additionally removes the bundled behavioral wedge, recovering the full no-behavioral baseline; the two are tracked separately because removing the wedge has a normatively contingent welfare interpretation (preference vs. bias).

Demand-side architecture. The bundled behavioral wedge admits a distinct policy lever that supply-side reforms cannot replicate: changing the choice architecture so that annuitization is the default rather than an opt-in decision. Removing the wedge under the production calibration (i.e., reporting the no-behavioral baseline directly without applying the multiplicative reduction) raises predicted ownership to 1.7%. Under bundled identification, this counterfactual is interpreted as removing the entire bundled wedge—the joint footprint of source-dependent utility, narrow framing, and choice-architecture salience that the UK reform reveals when compulsion lifts—rather than removing one named force. The counterfactual is well-defined as an aggregate wedge-removal intervention; sub-decomposing it into separate SDU / PED / salience contributions is not data-identified. The UK 2015 pension-freedoms reform—the calibration anchor for the bundle—produced a 75–87 percentage-point drop in DC-pot annuity ownership when regulatory compulsion was removed (the pre-reform regime enforced annuitization through a 55% tax penalty on alternative withdrawals; the post-reform regime is opt-in with marginal-rate tax neutrality) ([Association of British Insurers, 2020](#)); the [Goda et al. \(2014\)](#) federal Thrift Savings Plan annuity option records an ≈ 86 pp household-level default-vs-opt-in gap, a quasi-experimental signal that supports the order of magnitude of the bundled wedge effect. Demand-side architecture is quantitatively comparable to the largest supply-side reforms considered here, while requiring no change to insurer pricing or product design.

Table 13: Welfare Gain from Annuity Access Under Policy Counterfactuals (CEV, %)

Wealth	Baseline	Group pricing	Real annuity	Best feasible
<i>Panel A: Good Health, DFJ Bequests</i>				
\$50K	0.00	0.00	0.00	0.00
\$100K	0.00	0.00	0.00	0.00
\$200K	0.00	0.00	0.00	1.47
\$500K	0.24	0.75	0.42	4.43
\$1000K	0.86	1.53	1.07	5.52
<i>Panel B: Fair Health, DFJ Bequests</i>				
\$50K	0.00	0.00	0.00	0.00
\$100K	0.00	0.00	0.00	0.00
\$200K	0.00	0.00	0.00	0.80
\$500K	0.07	0.49	0.21	3.64
\$1000K	0.59	1.12	0.77	4.59
<i>Panel C: Good Health, No Bequests</i>				
\$50K	0.00	0.00	0.00	0.00
\$100K	0.00	0.00	0.00	0.00
\$200K	0.00	0.33	0.00	2.27
\$500K	3.20	4.68	3.32	8.40
\$1000K	7.72	9.81	7.93	14.27

CEV: consumption-equivalent variation (% of lifetime consumption agent would pay for annuity market access). Welfare model uses representative SS income (\$18,500/yr, median quartile). Baseline: MWR=0.87, 2% inflation, $\psi = 0.981$, $\psi_{\text{purchase}} = 0.0163$. Group pricing: MWR=0.90. Real annuity: 0% inflation, MWR=0.87. Best feasible combines group pricing, real annuity, no survival pessimism ($\psi = 1.0$), and default architecture ($\psi_{\text{purchase}} = 0$).

Welfare under reform. Under baseline pricing, the welfare gain from market access is small at every wealth-health cell—0.0% for Good health with DFJ bequests at \$200,000—because most calibrated households would not optimally annuitize at current prices. Policy interventions that change the optimal choice generate the larger welfare gains. Group pricing ($MWR = 0.90$) roughly quadruples CEV at moderate wealth: 0.0% at \$200,000 (Good health, DFJ). The “best feasible” scenario (group pricing combined with default architecture, which removes the bundled behavioral wedge) extends gains further: 0.0% at \$100,000 and 1.5% at \$200,000. CEV is near zero at high wealth under baseline, group-pricing-only, and real-annuity-only scenarios. The large high-wealth gains in the best-feasible column arise only when the default-architecture counterfactual removes the bundled wedge, so that column should be interpreted as a behavioral-choice-architecture lever rather than a price-only reform. The values reported in the welfare table are a CRRA value-ratio welfare index. Because the model includes non-homothetic bequests and a consumption floor, these values should be read as directional access-value statistics rather than exact compensating variation; the ranking and signs across scenarios are preserved by the approximation.

The welfare analysis reframes the annuity puzzle. For the median retiree under current market conditions and the conservative (ABI-anchored) bracket calibration, non-annuitization is not a mistake; it is the correct response to loaded pricing, nominal payment erosion, correlated health and survival risk, pre-existing Social Security coverage, declining consumption needs, and the at-purchase narrow-framing penalty. Consistent with that, the welfare gain from *access* alone is small. The welfare case for annuitization emerges only when policy reforms also change the optimal choice—supply-side cost reductions and demand-side default architecture together raise CEV at moderate wealth from near zero to roughly 1%–1.5% of consumption.

5 Robustness

Table 8 in the appendix reports the full sensitivity analysis. Key dimensions are summarized here.

5.1 Risk aversion

The sensitivity of predicted ownership to γ is documented in Table 8 and discussed in Section 4.7. At the baseline $\gamma = 2.5$, the six-channel rational model predicts 8.5%. Ownership is 0.0% for $\gamma \leq 2.0$, rises through 0.0% at $\gamma = 2.3$ and 0.1% at $\gamma = 2.4$, and reaches 18.1% at $\gamma = 3.0$. The eight-channel rational+preference model shifts these values downward, producing 1.8% at the baseline.

Inflation has competing effects in the nominal-annuity calibration: it erodes the real value of later payments, but it also raises the initial nominal payout through the insurer’s pricing, which uses the nominal discount rate. The net effect on demand depends on which side dominates at the calibrated parameters. At the baseline $\gamma = 2.5$, predicted ownership rises monotonically over the empirically defensible range, from 4.7% at $\pi = 1\%$ to 0.7% at $\pi = 3\%$, indicating that the higher initial nominal payout dominates payment erosion at the calibrated margin.

5.2 Health-mortality correlation

The hazard multipliers govern the strength of the Reichling–Smetters mechanism. The R-S functional estimates $([0.45, 1.0, 3.5])$ produce 3.2% ownership; the HRS self-reported health estimates $([0.57, 1.0, 2.7])$ produce 6.4%; the conservative specification $([0.60, 1.0, 2.0])$ produces 10.6%. An age-varying specification that linearly interpolates the HRS estimates across the three age bands in Section 3 produces 7.3%, reflecting the compression of the health-mortality gradient at older ages (the Poor-health multiplier falls from 3.29 at ages 65–74 to 1.82 at 85+). The qualitative pattern—wider health gradients reduce annuity

demand—holds across all parameterizations.

5.3 Money’s worth ratio

Table 14 reports ownership under MWR values from the policy counterfactual exercise.

Table 14: Predicted Ownership by Money’s Worth Ratio

MWR	Ownership (%)
0.82	0.0
0.85	0.4
0.87 (baseline)	1.7%
0.90	5.8
0.95	15.4

Production model at the no-behavioral baseline (rational + preferences + $\chi_{\text{LTC}} = 0.7$); the bundled behavioral wedge is applied as multiplicative post-processing ($\rho = 0.179$ at the production midpoint).

5.4 Survival pessimism

Varying ψ from 0.970 to 1.000 (objective beliefs) shifts predicted ownership from 5.1% to 21.4%. At the baseline $\psi = 0.960$, the model predicts 8.7%. The steep sensitivity confirms that survival pessimism is a quantitatively meaningful channel, consistent with its Shapley value of 5.9 pp.

5.5 Additional robustness checks

Grid convergence. At the 9-node production quadrature rule, predicted ownership ranges from 25.6% at the medium grid (60×20) to 23.6% at the fine grid (100×40), bracketing the production grid’s no-behavioral baseline (Table 9). The coarse grid (40×15) deviates more substantially (29.6%) and is reported only as a diagnostic. The 2 pp variation across

the production-and-finer range is small relative to the UK retention sensitivity range that drives the bundled-wedge bracket.

Quadrature convergence. Binary ownership stabilizes from 9-node Gauss–Hermite onward (Table 4, Panel B): values at $n_{\text{quad}} \in \{9, 11, 13, 15\}$ are 25.47%, 25.75%, 25.65%, and 25.47% at the diagnostic specification, a range of 0.28 pp. Lower-order rules ($n_{\text{quad}} \leq 7$) deviate by up to 2.7 pp because the heavy lognormal medical-expense tail interacts with the Medicaid floor. Nine-node quadrature is adopted as the smallest order with negligible numerical error relative to the channel contributions; Appendix G provides details.

Bequest specification. Under the production ten-channel model at the conservative bracket end, the DFJ luxury-good specification ($\theta = 56.96$, $\kappa = \$272,628$) produces 1.7% ownership; removing bequests raises this to 21.0%. See Section 2 for discussion of why θ and κ must be interpreted jointly.

Shapley robustness. The exact Shapley decomposition computes over all $2^{10} = 1,024$ channel subsets, providing an order-independent attribution. Among demand-suppressing channels, the narrow-framing purchase-event disutility (Force B) has the largest Shapley value (39.6 pp), followed by pricing loads (14.0 pp) and the combined Med+R-S health-mortality channel (8.6 pp). The two behavioral channels enter with opposite-sign Shapley values—source-dependent utility contributes a negative Shapley value (-13.4 pp, raising demand) while the purchase-event disutility contributes a positive one (39.6 pp, suppressing demand)—a pattern consistent with operation on distinct decision margins rather than redundant parameterizations of the same wedge.

Numerical accuracy. Euler equation residuals, grid convergence, and quadrature convergence diagnostics are reported in Appendices B, G, and K.

6 Discussion

This paper is a calibrated structural exercise, not a formal estimation. All parameters are drawn from published sources. The main contribution is therefore quantitative accounting: how far the standard rational channels go, what the added preference channels contribute, and which mechanisms matter most once they are evaluated together.

Under the baseline calibration ($\gamma = 2.5$, $\text{MWR} = 0.87$), the six standard rational channels predict 8.5% ownership relative to a frictionless population benchmark of 41.8%—substantially overshooting both empirical targets and echoing the residual overprediction in [Pashchenko \(2013\)](#). The two preference channels bring the prediction to 1.8%. The bundled behavioral wedge (Force A and Force B jointly identified from the UK 2015 reform) then brings predicted ownership to approximately 6% at the production midpoint, with a UK retention bracket of approximately [4.6%, 8.8%] and a wider sensitivity range across alternative anchors of [2.5%, 8.8%]. Both HRS measures of US lifetime annuity ownership—the conventional any-annuity income proxy (3.34%, 95% CI [2.89%, 3.85%]) and the cleaner lifetime-contract indicator (2.02%, 95% CI [1.68%, 2.43%])—fall within the wider sensitivity range. The paper is therefore strongest when read as a layered result: rational and preference channels alone substantially overshoot, and the bundled behavioral wedge—calibrated from a single external moment that activated multiple behavioral channels at once—closes the residual gap with US empirical observation at the production midpoint while introducing the policy-relevant default-vs-opt-in margin.

The exact Shapley decomposition provides the clearest attribution hierarchy. Among demand-suppressing channels, the narrow-framing purchase penalty (Force B) is the largest, followed by pricing loads and the Reichling–Smetters health-mortality correlation; age-varying consumption needs, survival pessimism, and medical expenditure risk form the next tier. Source-dependent utility (Force A) and Social Security enter as boosters with negative Shapley values. State-dependent utility and bequest motives are empirically plausible but quantitatively negligible in this framework. That ranking is more informative than a single

overall ownership match.

The Social Security results illustrate the complement-substitute duality. In the frictionless environment, SS crowds *in* private annuity demand by providing an income floor that makes annuitization feasible. Once pricing loads and inflation erosion are active, the same SS income crowds *out* private annuity demand. The resulting non-monotonic response to SS cuts is a substantive public-finance result, not just a robustness check.

The baseline MWR of 0.87 is drawn from [Mitchell et al. \(1999\)](#). Recent evidence suggests that annuity pricing has improved, with population-table MWRs of approximately 0.87 for age-65 immediate annuities in 2024 market conditions. At $\text{MWR} = 0.85$, the model predicts 0.4% ownership (Table 14), well above observed rates. This sensitivity is not a weakness so much as a central substantive finding: pricing remains the highest-leverage policy margin in this market.

Several modeling choices condition the results.

Marital structure. The model treats all retirees as single. [Kotlikoff and Spivak \(1981\)](#) showed that marriage provides partial longevity insurance through intra-household risk sharing, which would further reduce annuity demand among married couples.

Housing wealth. Housing wealth is excluded. [Pashchenko \(2013\)](#) found that housing illiquidity reduces annuity demand by constraining the liquid wealth available for annuitization.

Subjective survival belief structure. The survival pessimism channel uses a simple proportional scaling of one-year survival rates; the actual structure of subjective belief distortions may be more complex ([O’Dea and Sturrock, 2023](#)).

Age-varying consumption needs heterogeneity. The age-varying consumption needs channel assumes a uniform 2% annual decline; heterogeneity in this rate across health states and wealth levels is not modeled.

The behavioral content of the model lives in a jointly identified bundle containing three named forces: source-dependent utility (Force A), narrow-framing purchase-event disutility

(Force B), and choice-architecture salience (Force C). The bundle is not parameterized inside the Bellman equation; it is identified externally from UK 2015 and applied as multiplicative post-processing on the no-behavioral baseline. The Blanchett–Finke spending differential is descriptive evidence that SDU exists as a behavioral phenomenon in retirement spending data; it is not used as an identifying moment here. The companion paper on Social Security claiming (Tharp, 2026, under review) parameterizes SDU separately under a different identifying strategy. Both HRS measures of US lifetime annuity ownership are reported as out-of-sample comparisons against the model’s bundled prediction. The Goda et al. (2014) federal Thrift Savings Plan default evidence is qualitatively consistent with the model’s implied default-architecture elasticity (which under joint identification reflects the combined effect of all three forces, not Force B alone), providing supporting context rather than a clean external validation. Hu and Scott (2007) showed analytically that prospect theory loss aversion reduces optimal annuitization substantially; the present paper operationalizes that mechanism within the broader bundled wedge, with parameter discipline from UK post-reform microdata. Ameriks et al. (2020) document the importance of long-term care aversion in late-life saving decisions (operationalized in this paper through the public-care aversion channel χ_{LTC}). Madrian and Shea (2001) showed that defaults dominate retirement savings behavior more generally; in this paper that finding is captured by Force C within the bundle, operationalized through the demand-side counterfactual rather than as a separately identified channel.

Deferred income annuities (DIAs) and qualified longevity annuity contracts (QLACs) have received attention as alternatives to immediate annuities. In an extension (Appendix E), DIA ownership rates are similar to SPIA rates. The lower MWR on deferred products—Wettstein et al. (2021) estimate MWRs of 0.50 for DIA-80 and 0.45 for DIA-85—offsets the actuarial advantage of concentrating payments in late life.

For policy, the counterfactual analysis identifies two distinct levers operating on different margins. Supply-side reform (group pricing at $MWR = 0.90$) raises the value of annuitiza-

tion. Demand-side reform (annuitization as default) is interpreted as removing the bundled wedge—reporting the no-behavioral baseline directly without the multiplicative reduction. Removing the wedge raises predicted ownership from the production midpoint to 1.7%, an elasticity in the range of the UK 2015 quasi-experiment and the federal TSP default evidence. These levers are not substitutes: pricing reform without default architecture leaves substantial residual demand suppression from the bundled behavioral wedge, and default architecture without pricing reform fails to capture households for whom the rational valuation gap is binding. The welfare case for annuitization is real for wealthy households even at current pricing, and becomes substantial under reform.

7 Conclusion

This paper built a lifecycle model that begins from a standard rational and preference-based account of annuity demand and then adds a *bundled behavioral wedge* containing three named channels (Force A: source-dependent utility; Force B: narrow-framing purchase-event disutility; Force C: choice-architecture salience) jointly identified from a single external moment, the UK 2015 pension-freedoms reform. Without the bundle, the model substantially overshoots both empirical targets; with it, the production calibration predicts approximately 6% US voluntary ownership at the production midpoint, with a UK retention bracket of [4.6%, 8.8%] and a wider sensitivity range across alternative anchors of [2.5%, 8.8%]. Both HRS measures of US lifetime annuity ownership—the conventional any-annuity income proxy (3.34%, 95% CI [2.89%, 3.85%]) and the cleaner lifetime-contract indicator (2.02%, 95% CI [1.68%, 2.43%])—fall within the wider sensitivity range.

Results have a three-layer structure. Six standard rational channels predict 8.5% ownership relative to a frictionless benchmark of 41.8%. Two preference channels—age-varying consumption needs and state-dependent utility—bring this to 1.8%. Adding the structural public-care aversion channel (χ_{LTC}) yields the no-behavioral baseline of 1.7%. The

bundled behavioral wedge is then applied to this baseline as multiplicative post-processing under proportional-wedge transport: predicted US ownership equals the no-behavioral baseline times the UK retention factor $\rho = 0.179$ at the production midpoint, yielding 7.5%. The bundle is not parameterized inside the model; the wedge is implemented as a deterministic scalar reduction. The decomposition into named components (SDU, PED, choice-architecture salience) is conceptual; a single bundling moment cannot separately attribute the swing to any one component. Each layer is a substantive contribution, and the calibration is disciplined out of sample at every stage: no US moment is used to set the bundle.

Exact Shapley values, computed over all $2^9 = 512$ subsets of the nine model-internal channels, provide the paper’s preferred attribution for the no-behavioral baseline. Among demand-suppressing channels, pricing loads and the combined Med+R-S health-mortality channel are the two largest contributors; age-varying needs, survival pessimism, and inflation erosion form the next tier. Social Security enters with a negative Shapley value, raising annuity demand in frictionless environments but crowding out private annuitization once market frictions are active. The bundled behavioral wedge does not appear in the Shapley enumeration: it is applied as multiplicative post-processing on ownership rates, separate from the channel-by-channel attribution.

Welfare and policy implications follow from the same calibration. Under current pricing, the welfare gain from annuity market access alone is small (under DFJ bequests, 0.0% at \$200,000 in Good health), because most calibrated households would not optimally buy at current prices. Larger welfare gains require policy reforms that change the optimal choice: group pricing roughly quadruples CEV at moderate wealth, and the best-feasible welfare package (group pricing combined with default architecture, which removes the bundled behavioral wedge) reaches 1.5% at \$200,000. The welfare interpretation of removing the bundled wedge is normatively contingent on treating the underlying SDU/PED/salience phenomena as preference rather than bias. Two distinct policy levers operate on different margins: supply-side reforms (group pricing) raise the value of annuitization; demand-side

reforms (default architecture, removing the bundled wedge) operate on the wedge magnitude.

Research framing should therefore shift from “why don’t retirees annuitize?” to “which channels matter most, and what mix of supply-side and demand-side interventions would make annuitization welfare-improving, and for whom?”

Data Availability

Replication code (Julia) and processed calibration data are available at <https://github.com/DerekTharp/annuity-puzzle-model>. Health transition matrices and wealth distributions are computed from the RAND HRS Longitudinal File (public use, available at <https://hrsdata.isr.umich.edu>). All other calibration parameters are taken from published sources cited in the text.

References

- Aguiar, Mark and Erik Hurst**, “Deconstructing Life Cycle Expenditure,” *Journal of Political Economy*, 2013, 121 (3), 437–492.
- Ameriks, John, Joseph Briggs, Andrew Caplin, Matthew D. Shapiro, and Christopher Tonetti**, “Long-Term-Care Utility and Late-in-Life Saving,” *Journal of Political Economy*, 2020, 128 (6), 2375–2451.
- Association of British Insurers**, “Pension Freedoms: Five Years On,” ABI Industry Report. Reports annuity sales declines from approximately 353,000 in 2013 to 60,000–90,000 annually post-2015 reform. 2020.
- Barberis, Nicholas and Ming Huang**, “Preferences with Frames: A New Utility Specification that Allows for the Framing of Risks,” *Journal of Economic Dynamics and Control*, 2009, 33 (8), 1555–1576.

- Bernheim, B. Douglas and Antonio Rangel**, “Beyond Revealed Preference: Choice-Theoretic Foundations for Behavioral Welfare Economics,” *Quarterly Journal of Economics*, 2009, 124 (1), 51–104.
- Beshears, John, James J. Choi, David Laibson, and Brigitte C. Madrian**, “How Are Preferences Revealed?,” *Journal of Public Economics*, 2008, 92 (8–9), 1787–1794.
- Blanchett, David and Michael S. Finke**, “Guaranteed Income and the Retirement Spending Puzzle,” *Working Paper*, 2024. Documents that retirees consume approximately 80% of guaranteed income but only 50% of common withdrawal benchmarks (e.g., the 4% rule) from portfolio wealth.
- and —, “Guaranteed Income: A License to Spend,” *The Journal of Retirement*, 2025. Forthcoming. Documents that retirees with annuitized income spend at higher rates than equivalent-wealth retirees drawing from portfolios, implying an asymmetric valuation of annuitized vs. liquid wealth.
- Brown, Jeffrey R., Jeffrey R. Kling, Sendhil Mullainathan, and Marian V. Wrobel**, “Why Don’t People Insure Late-Life Consumption? A Framing Explanation of the Under-Annuitization Puzzle,” *American Economic Review: Papers and Proceedings*, 2008, 98 (2), 304–309.
- Cannon, Edmund, Ralph Stevens, and Ian Tonks**, “Price Efficiency in the Dutch Annuity Market,” *Journal of Pension Economics and Finance*, 2016, 15 (1), 1–26. Provides European context for UK annuity-market structure including tax-driven retention pre-2015.
- Chalmers, John and Jonathan Reuter**, “Is Conflicted Investment Advice Better than No Advice?,” Working Paper 18158, National Bureau of Economic Research June 2012.
- Chetty, Raj**, “A New Method of Estimating Risk Aversion,” *American Economic Review*, 2006, 96 (5), 1821–1834.

- Cribb, Jonathan and Carl Emmerson**, “What Happens When Employers Are Obligated to Nudge? Automatic Enrolment and Pension Saving in the UK,” *NBER Working Paper*, 2018, 25697. Documents the structure and behavioral response to UK pension reforms.
- Davidoff, Thomas, Jeffrey R. Brown, and Peter A. Diamond**, “Annuities and Individual Welfare,” *American Economic Review*, 2005, 95 (5), 1573–1590.
- Dushi, Irena and Anthony Webb**, “Household Annuitization Decisions: Simulations and Empirical Analyses,” *Journal of Pension Economics and Finance*, 2004, 3 (2), 109–143.
- Financial Conduct Authority**, “Retirement income market data 2024/25,” FCA published statistics. Annuity share of DC pension pots accessed for the first time stood at approximately 9% in 2024/25 despite gilt yields at 14-year highs, indicating broadly stable post-2015 retention regime across a full rate cycle and a small rate-environment confound for ψ_{purchase} identification. 2025.
- Finkelstein, Amy, Erzo F.P. Luttmer, and Matthew J. Notowidigdo**, “What Good Is Wealth Without Health? The Effect of Health on the Marginal Utility of Consumption,” *Journal of the European Economic Association*, 2013, 11 (S1), 221–258.
- Goda, Gopi Shah, Colleen Flaherty Manchester, and Aaron J. Sojourner**, “What Will My Account Really Be Worth? Experimental Evidence on How Retirement Income Projections Affect Saving,” *Journal of Public Economics*, 2014, 119, 80–92. Documents default-vs-opt-in elasticity in federal Thrift Savings Plan annuity option (≈ 86 pp gap).
- Hastings, Justine S. and Jesse M. Shapiro**, “Fungibility and Consumer Choice: Evidence from Commodity Price Shocks,” *Quarterly Journal of Economics*, 2013, 128 (4), 1449–1498.
- Hu, Wei-Yin and Jason S. Scott**, “Behavioral Obstacles in the Annuity Market,” *Financial Analysts Journal*, 2007, 63 (6), 71–82.

- James, Estelle and Xue Song**, “Annuities Markets Around the World: Money’s Worth and Risk Intermediation,” CeRP Working Paper 16/01, Center for Research on Pensions and Welfare Policies 2006.
- Jones, John Bailey, Mariacristina De Nardi, Eric French, Rory McGee, and Justin Kirschner**, “The Lifetime Medical Spending of Retirees,” Working Paper 24599, National Bureau of Economic Research May 2018.
- Kotlikoff, Laurence J. and Avia Spivak**, “The Family as an Incomplete Annuities Market,” *Journal of Political Economy*, 1981, 89 (2), 372–391.
- Lockwood, Lee M.**, “Bequest Motives and the Annuity Puzzle,” *Review of Economic Dynamics*, 2012, 15 (2), 226–243.
- Madrian, Brigitte C. and Dennis F. Shea**, “The Power of Suggestion: Inertia in 401(k) Participation and Savings Behavior,” *The Quarterly Journal of Economics*, 2001, 116 (4), 1149–1187.
- Mitchell, Olivia S., James M. Poterba, Mark J. Warshawsky, and Jeffrey R. Brown**, “New Evidence on the Money’s Worth of Individual Annuities,” *American Economic Review*, 1999, 89 (5), 1299–1318.
- Nardi, Mariacristina De**, “Wealth Inequality and Intergenerational Links,” *Review of Economic Studies*, 2004, 71 (3), 743–768.
- , **Eric French, and John B. Jones**, “Why Do the Elderly Save? The Role of Medical Expenses,” *Journal of Political Economy*, 2010, 118 (1), 39–75.
- O’Dea, Cormac and David Sturrock**, “Survival Pessimism and the Demand for Annuities,” *The Review of Economics and Statistics*, 2023, 105 (2), 442–457.
- Pashchenko, Svetlana**, “Accounting for Non-Annuitization,” *Journal of Public Economics*, 2013, 98, 53–67.

- Peijnenburg, Kim, Theo Nijman, and Bas J. M. Werker**, “The Annuity Puzzle Remains a Puzzle,” *Journal of Economic Dynamics and Control*, 2016, 70, 18–35.
- Poterba, James, Steven Venti, and David Wise**, “The Composition and Drawdown of Wealth in Retirement,” *Journal of Economic Perspectives*, 2011, 25 (4), 95–118.
- Reichling, Felix and Kent Smetters**, “Optimal Annuitization with Stochastic Mortality and Correlated Medical Costs,” *American Economic Review*, 2015, 105 (11), 3273–3320.
- Shefrin, Hersch M. and Richard H. Thaler**, “The Behavioral Life-Cycle Hypothesis,” *Economic Inquiry*, 1988, 26 (4), 609–643.
- Thaler, Richard H.**, “Mental Accounting Matters,” *Journal of Behavioral Decision Making*, 1999, 12 (3), 183–206.
- Tharp, Derek**, “Revisiting the Social Security Claiming Puzzle: Behavioral Preferences as Rational Explanations for Early Claiming,” Under review (revise and resubmit) at *Financial Planning Review*. Calibrates source-dependent utility parameter $\lambda_W = 0.625$ to the same Blanchett-Finke spending differential used for source-dependent utility in the present paper. 2026.
- Tversky, Amos and Daniel Kahneman**, “Advances in Prospect Theory: Cumulative Representation of Uncertainty,” *Journal of Risk and Uncertainty*, 1992, 5 (4), 297–323.
- Wettstein, Gal, Alicia H. Munnell, Wenliang Hou, and Nilufer Gok**, “The Value of Annuities,” Working Paper CRR WP 2021-5, Center for Retirement Research at Boston College March 2021.
- Yaari, Menahem E.**, “Uncertain Lifetime, Life Insurance, and the Theory of the Consumer,” *The Review of Economic Studies*, 1965, 32 (2), 137–150.




Figure 1: Sequential decomposition of predicted annuity ownership (six-channel rational model). Each bar shows predicted ownership after adding the labeled channel to all previous channels. The dashed line marks observed ownership in the present sample (3.34% on the conventional any-annuity income proxy, 2.02% on the cleaner lifetime contract indicator).




Figure 2: Predicted ownership by coefficient of relative risk aversion (γ). All other parameters at baseline values ($\pi = 2\%$, MWR = 0.87, DFJ bequests). The shaded band marks observed ownership of 3–6%.




Figure 3: Predicted ownership by hazard multiplier specification. Each specification varies the Good- and Poor-health multipliers while holding Fair at 1.0. Labels indicate the empirical source for each calibration. The shaded band marks observed ownership of 3–6%.




Figure 4: Optimal annuitization fraction (α^*) at age 65 by initial wealth, for Good health ($H = 1$). Left panel: no bequest motive, pricing loads and inflation active. Right panel: full model (DFJ bequests, medical risk, Reichling–Smetters correlation, pricing loads, inflation). Both panels use $MWR = 0.87$ and $\pi = 2\%$.



Figure 5: Consumption-equivalent variation (%) from annuity market access, by initial wealth and health status. No-bequest specification ($\theta = 0$). Full model with medical costs, Reichling–Smetters correlation, $MWR = 0.87$, and 2% inflation.